

# **Operating Systems with Linux**

**John O’Gorman**

palgrave



# Figures

## **1 Introduction**

- 1.1 Overview of a computer system
- 1.2 Interfaces to an operating system
- 1.3 Layered structure
- 1.4 Poorly designed system
- 1.5 Modular operating system
- 1.6 Microkernel with two personalities

## **2 Interfaces to an operating system**

- 2.1 Memory layout for global variables
- 2.2 Checking error values
- 2.3 Indexing into the system call table
- 2.4 Calling a system service
- 2.5 Computer memory
- 2.6 Dual-ported memory
- 2.7 Keyboard interface with computer
- 2.8 Interface using interrupt line
- 2.9 Interrupt arbitration by controller
- 2.10 Direct memory access

## **3 Process manager**

- 3.1 Program to illustrate change of state
- 3.2 States as program executes
- 3.3 Overlap of processor and devices

- 3.4 An underutilised system
- 3.5 High utilisation of CPU and disk drives
- 3.6 How the system keeps track of processes
- 3.7 Data structure representing a process
- 3.8 Program to illustrate fork()
- 3.9 Function call and thread creation
- 3.10 An entry in a wait queue
- 3.11 Three processes waiting for a resource
- 3.12 Process states and transitions
- 3.13 Running in user or kernel mode
- 3.14 The volatile environment of a process
- 3.15 Data saved by context switcher
- 3.16 Context switching
- 3.17 Processes on the ready and wait queues
- 3.18 Two different priority queues

## **4 Concurrency**

- 4.1 Code to increment counter
- 4.2 Code to decrement counter
- 4.3 General structure of a cooperating program
- 4.4 First attempt at solution for two processes
- 4.5 Compiled version of entry protocol
- 4.6 Algorithms using turn
- 4.7 Algorithm using two flags
- 4.8 Setting the flag before testing
- 4.9 Algorithm for mutual exclusion
- 4.10 Entry protocols for both processes
- 4.11 Eight contending processes
- 4.12 Eisenberg and McGuire algorithm
- 4.13 Lamport’s algorithm

## **5 Low level IPC mechanisms**

- 5.1 A selection of signal values
- 5.2 Changing the signal mask
- 5.3 Data structure tracking signal handlers
- 5.4 Entry protocol using test and set
- 5.5 Entry protocol using exchange
- 5.6 Semaphore for mutual exclusion

- 5.7 Successive states of a mutual exclusion semaphore
- 5.8 Semaphore for synchronisation
- 5.9 An alternative for synchronisation
- 5.10 Semaphore for resource allocation
- 5.11 Circular buffer with eight slots
- 5.12 Algorithm for producer
- 5.13 Algorithm for consumer
- 5.14 Algorithm for one of many producers
- 5.15 Algorithm for one of many consumers
- 5.16 Entry protocol for a reader
- 5.17 Exit protocol for a reader
- 5.18 Solution 2
- 5.19 Solution 3
- 5.20 Reader prologue code
- 5.21 Incorrect reader prologue
- 5.22 Reader epilogue code
- 5.23 Writer prologue code
- 5.24 Writer epilogue code
- 5.25 Solution 4
- 5.26 Reader epilogue code
- 5.27 Writer prologue code
- 5.28 Writer epilogue code
- 5.29 Timing relationship of locks and semaphores
- 5.30 Data structure representing a semaphore set
- 5.31 Data structure controlling access to a semaphore
- 5.32 Data structure representing an individual semaphore
- 5.33 System V semaphore layout
- 5.34 Data structure representing a waiting process
- 5.35 Processes waiting on the semaphore set
- 5.36 Eventcount for synchronisation
- 5.37 Eventcount and sequencer for mutual exclusion
- 5.38 Eventcounts with one producer, one consumer

## **6 Higher level mechanisms for IPC**

- 6.1 Outline of message passing
- 6.2 A message queue
- 6.3 Data structure representing a message queue
- 6.4 Structure of an actual message
- 6.5 System V message queue layout

- 6.6 Data structure representing a message
- 6.7 Readers/writers using conditional critical regions
- 6.8 Implementation of CWAIT
- 6.9 Implementation of CSIGNAL
- 6.10 An illustration of a monitor
- 6.11 Monitor to allocate a single resource
- 6.12 Monitor to manage a bounded buffer
- 6.13 Monitor to implement reader/writer interlock

## **7 Deadlock**

- 7.1 A resource allocation graph
- 7.2 A resource allocation graph for four processes
- 7.3 Safe, unsafe, and deadlocked system
- 7.4 A claims graph
- 7.5 One resource allocated
- 7.6 The banker’s algorithm
- 7.7 Safety algorithm
- 7.8 State of the system before request
- 7.9 Application of the safety algorithm
- 7.10 State of the system reflecting the request
- 7.11 Resource allocation graph
- 7.12 Wait-for graph
- 7.13 Algorithm to detect deadlock
- 7.14 A system which is not deadlocked
- 7.15 Application of the deadlock detection algorithm
- 7.16 A deadlocked system

## **8 Memory manager**

- 8.1 Hardware memory
- 8.2 The memory pyramid
- 8.3 Program as loaded by relocating loader
- 8.4 Virtual memory
- 8.5 Address modification in CPU
- 8.6 Protection and address modification
- 8.7 Compaction
- 8.8 A logical address
- 8.9 Address translation with segment table
- 8.10 A segment table entry

- 8.11 Searching a segmentation cache
- 8.12 Segmentation algorithm
- 8.13 Data structure representing a physical page
- 8.14 Tracking free physical memory
- 8.15 Linux page table structure
- 8.16 Format of page table entry
- 8.17 Address translation with cache and page table
- 8.18 Program to initialise an array
- 8.19 The root of the memory management data structures
- 8.20 Data structure controlling a region
- 8.21 Memory management data structures for a process
- 8.22 Possible values for the flags field
- 8.23 Operations on a region
- 8.24 Pages actually in memory
- 8.25 Example mapping of three regions
- 8.26 Relationship between page frames and page faults
- 8.27 Algorithm for NRU replacement
- 8.28 Effect of thrashing on CPU utilisation
- 8.29 Effect of window size on working set
- 8.30 An address interpreted as segment, page, offset
- 8.31 Algorithm for paged segmentation
- 8.32 Address translation with paged segmentation

## **9 Input and output**

- 9.1 Overview of I/O processing
- 9.2 A tree-structured directory
- 9.3 Structure of a directory entry
- 9.4 A directory with links
- 9.5 How a hard link is implemented
- 9.6 How a symbolic link is implemented
- 9.7 A directory structure with cycles
- 9.8 Inode for an open file
- 9.9 Inodes on LRU and hash lists
- 9.10 Mode bit values
- 9.11 Some functions of the inode interface
- 9.12 An element in the cache of recent lookups
- 9.13 Directory name lookup cache
- 9.14 A protection bitmap
- 9.15 Definition of file descriptor table

- 9.16 Descriptor and open file data structure
- 9.17 Structure representing an open stream
- 9.18 File opened by one user
- 9.19 Operations on open files
- 9.20 Algorithm for open
- 9.21 File opened by two users
- 9.22 Generic I/O algorithm
- 9.23 Generic algorithm for Unix I/O procedure
- 9.24 Data structure for polling a descriptor
- 9.25 Several processes waiting on a device
- 9.26 Structure representing a lock on a file
- 9.27 A lock request which blocks
- 9.28 How locks are recorded
- 9.29 A buffered input procedure
- 9.30 A buffer head
- 9.31 Bits representing the state of a buffer
- 9.32 Buffer cache

## **10 Regular file systems**

- 10.1 An entry in the linked list of file systems
- 10.2 Layout of a disk
- 10.3 Contiguous allocation
- 10.4 Linked allocation
- 10.5 Example of a file map
- 10.6 Layout of a partition with Ext2
- 10.7 A block group on disk
- 10.8 An Ext2 super block
- 10.9 Structure of a block group descriptor
- 10.10 Structure of an Ext2 disk inode
- 10.11 Allocation information in an inode
- 10.12 Single indirect block
- 10.13 Double indirect block
- 10.14 Fragments of two files in a block
- 10.15 An Ext2 directory entry
- 10.16 File systems as they exist on disk
- 10.17 Directory structure after attaching file system B
- 10.18 Data structure representing a mounted file system
- 10.19 Operations on a file system
- 10.20 Superblock for an Ext2 file system



- 10.21 Data structures involved in mounting a file system
- 10.22 Ext2 file system inode data in memory

## 11 Special files

- 11.1 Calling a Unix style driver
- 11.2 Device driver as a process
- 11.3 The character and block device switches
- 11.4 An entry in the character device switch
- 11.5 Algorithm for reading from a character device
- 11.6 Request block for block devices
- 11.7 Algorithm for strategy routine
- 11.8 Algorithm for interrupt service routine
- 11.9 Algorithm for writing to a character device
- 11.10 A basic view of a stream
- 11.11 Data structures constituting a minimal stream
- 11.12 A many-to-one multiplexor
- 11.13 A one-to-many multiplexor
- 11.14 A many-to-many multiplexor
- 11.15 Networking multiplexor
- 11.16 Data structure controlling terminal characteristics
- 11.17 Logging on using a pseudo terminal

## 12 IPC files

- 12.1 Setting up a pipe for interprocess communication
- 12.2 Data structure representing a pipe
- 12.3 Representation of a pipe
- 12.4 Supported address families
- 12.5 Data structure representing a protocol family
- 12.6 Domains registered in the kernel
- 12.7 Data structure representing a socket
- 12.8 State values for a socket
- 12.9 Data structure representing a protocol
- 12.10 Data structures after socket is allocated
- 12.11 Generic format of a socket address
- 12.12 Message header
- 12.13 Domain-specific control block
- 12.14 Hash chain of Unix domain control blocks
- 12.15 A protocol specific structure

- 12.16 Links for a Unix specific socket
- 12.17 Structure containing address of a Unix domain socket
- 12.18 A pair of connected sockets
- 12.19 Buffer used for communication over sockets
- 12.20 Data part of an sk\_buff

## **13 Distributed systems**

- 13.1 Networked computers
- 13.2 Workstations connected on a LAN
- 13.3 Diskless workstations on a LAN
- 13.4 Processor pool model
- 13.5 Combined model
- 13.6 Part of the Internet name space
- 13.7 Using a name server
- 13.8 Translating an Internet name
- 13.9 Overview of CORBA
- 13.10 The distributed computing environment

## **14 Communication**

- 14.1 Data structure describing a network interface
- 14.2 An Internet protocol control block
- 14.3 Internet socket address structure
- 14.4 Connecting two sockets
- 14.5 A physical frame
- 14.6 An ethernet header
- 14.7 Structure identifying a packet type
- 14.8 An IP header
- 14.9 A UDP header
- 14.10 A TCP header
- 14.11 Function call/return
- 14.12 Call/return with RPC
- 14.13 Format of RPC specification language
- 14.14 Interface definition for MATH program
- 14.15 Unix style credentials
- 14.16 Overview of RPC system
- 14.17 Structure of an RPC message
- 14.18 Body of a call message
- 14.19 Body of a reply message

## **15 Concurrency in distributed systems**

- 15.1 Events in process X
- 15.2 Events in process Y
- 15.3 Events in process Z
- 15.4 Centralised control of mutual exclusion
- 15.5 Distributed control of mutual exclusion
- 15.6 Successful two phase commit
- 15.7 Aborted two phase commit
- 15.8 A cascaded abort
- 15.9 A lock table
- 15.10 Algorithm for lock
- 15.11 Algorithm for unlock
- 15.12 Local and global wait-for graphs


## **16 Distributed file systems**

- 16.1 Situation before attaching remote files
- 16.2 Local file system after attaching remote files
- 16.3 A client and two servers
- 16.4 File system after mounting Dir4
- 16.5 File system after mounting Dir7
- 16.6 Mounting a remote file system
- 16.7 NFS specific super block information
- 16.8 Interaction between client and server
- 16.9 NFS specific inode information
- 16.10 Handling local and remote file systems

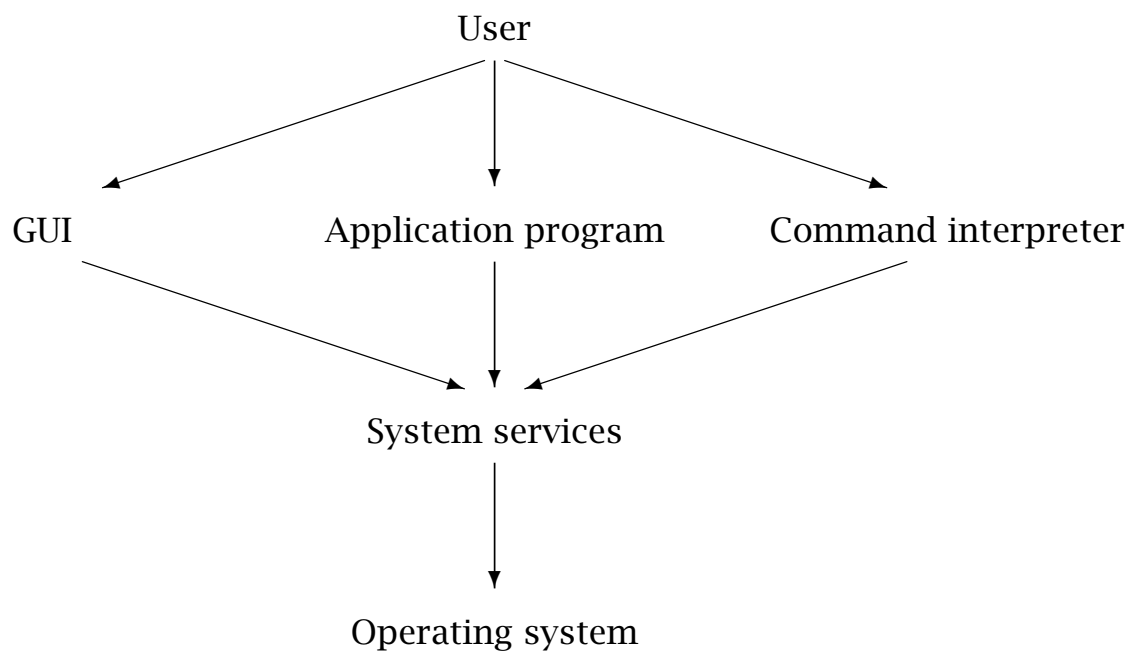
## **17 Fault tolerance and security**

- 17.1 Example access matrix
- 17.2 Sequence of messages with secret keys
- 17.3 Digital signature with secret key
- 17.4 Encryption for both security and authentication
- 17.5 Three phases in the Kerberos system
- 17.6 Initial exchange of messages with authentication server
- 17.7 Authenticating a user to a file server
- 17.8 Authenticating a user to a file server
- 17.9 Requesting service of a file server






**Figure 1.1:** Overview of a computer system



**Figure 1.2:** Interfaces to an operating system



**Figure 1.3:** Layered structure

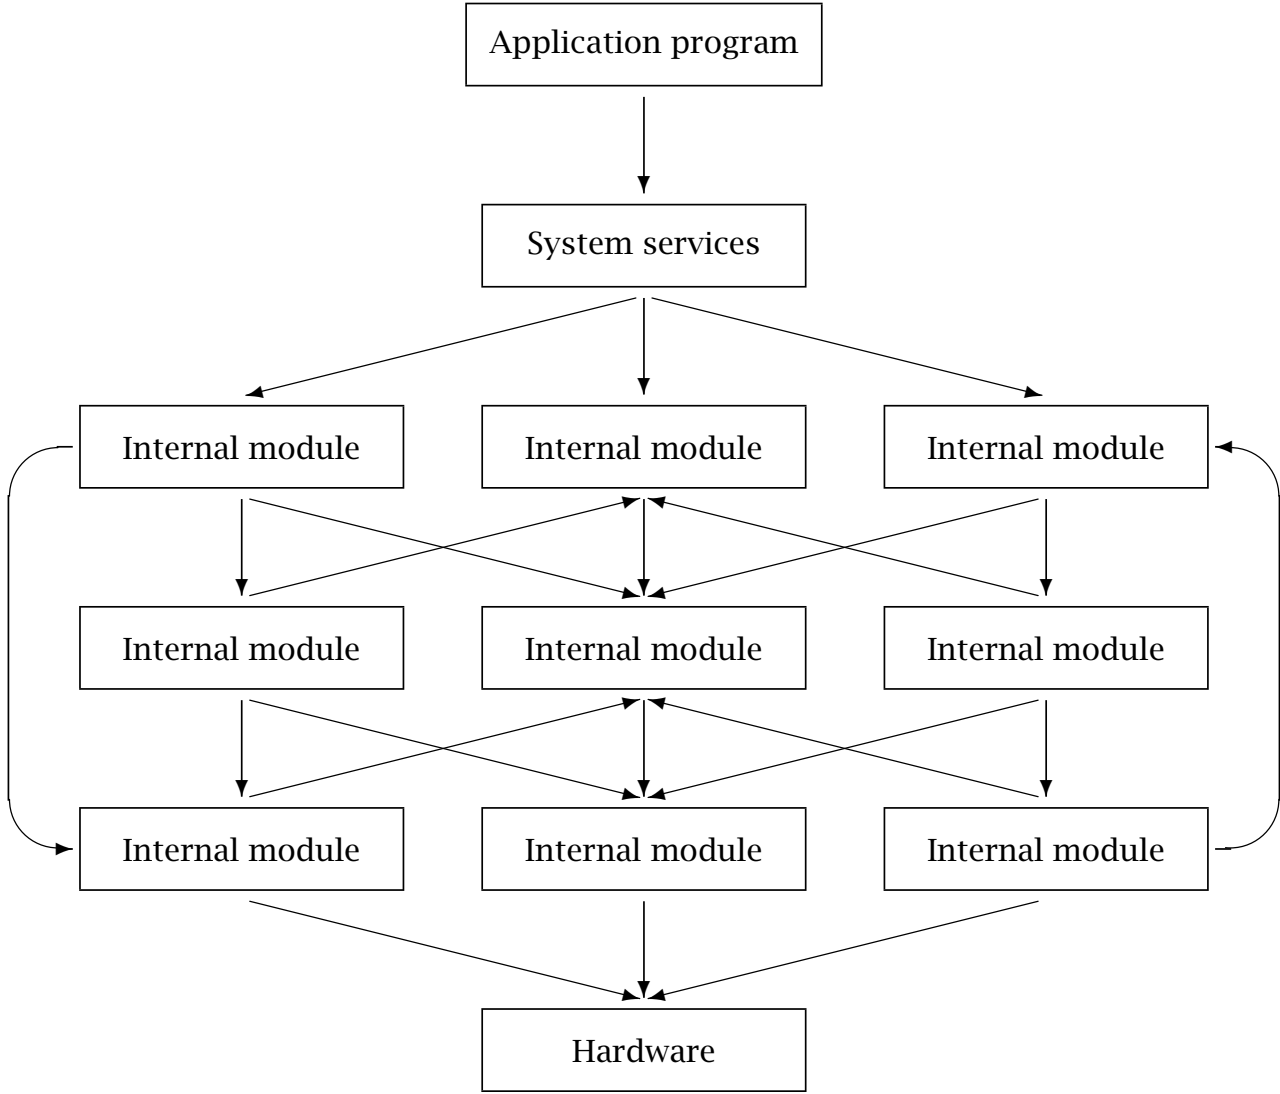


Figure 1.4: Poorly designed system



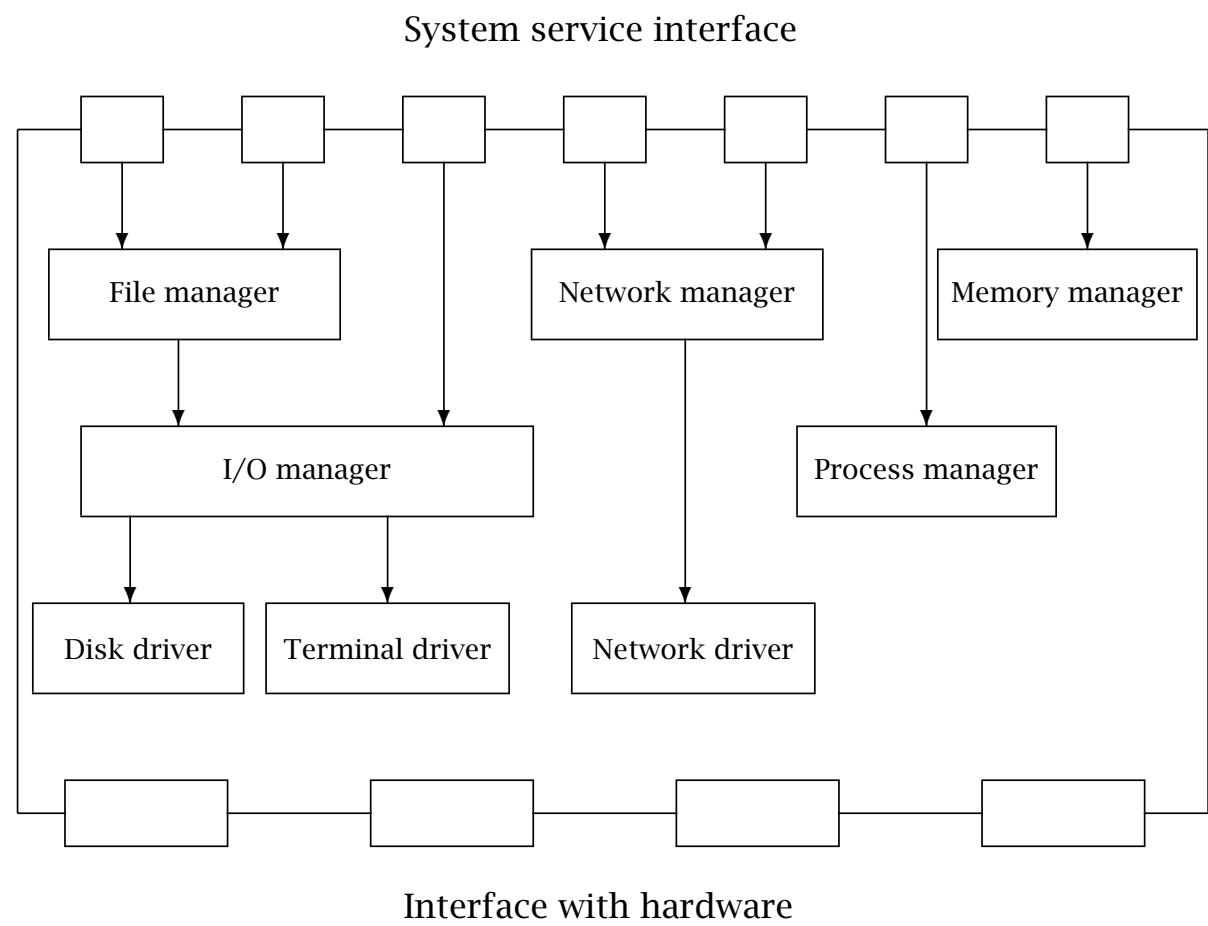


Figure 1.5: Modular operating system


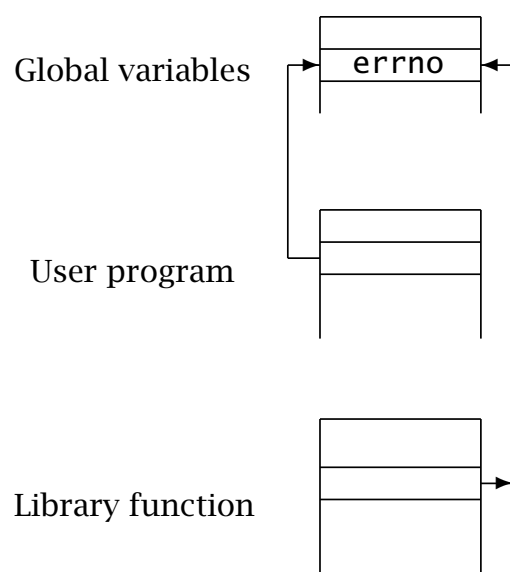


Figure 1.6: Microkernel with two personalities



**Figure 2.1:** Memory layout for global variables

```
#include <errno.h>

ret_val = syservice();
if (ret_val == -1) switch(errno)
    {case EAGAIN:  printf("one message");
      break;
     case EBADPARM: printf("another message");
    };
```

**Figure 2.2:** Checking error values







Figure 2.3: Indexing into the system call table



**Figure 2.4:** Calling a system service



**Figure 2.5:** Computer memory



**Figure 2.6:** Dual-ported memory



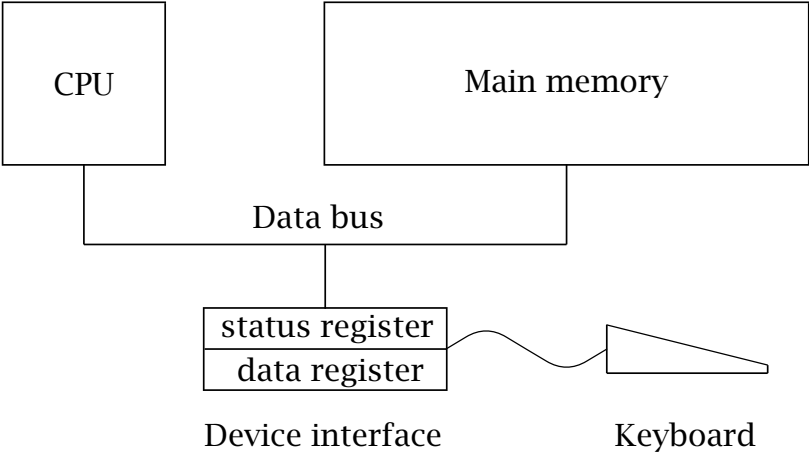




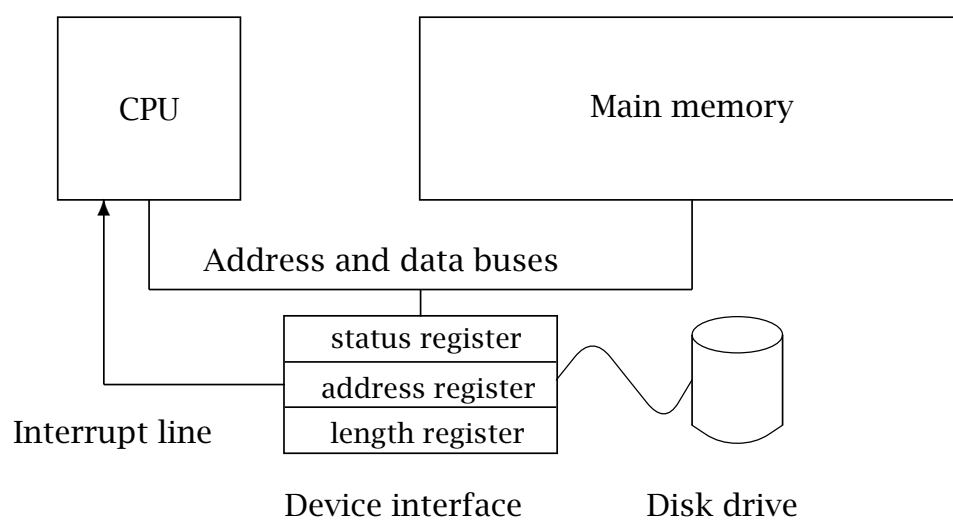
Figure 2.7: Keyboard interface with computer



**Figure 2.8:** Interface using interrupt line



**Figure 2.9:** Interrupt arbitration by controller




**Figure 2.10:** Direct memory access

```
1: MOV EAX, 7      ; load the value 7 into register EAX
2: INC EAX          ; increment register EAX
3: MOV total, EAX; store the value from register EAX to total
```

**Figure 3.1:** Program to illustrate change of state

	EIP	EAX	total
Initial state	1	0	0
State after 1st instruction	2	7	0
State after 2nd instruction	3	8	0
State after 3rd instruction	4	8	8

**Figure 3.2:** States as program executes






Figure 3.4: An underutilised system






Figure 3.5: High utilisation of CPU and disk drives

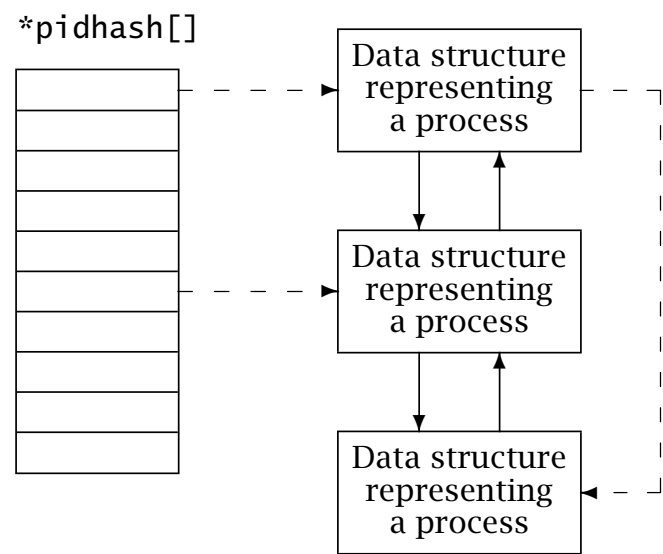


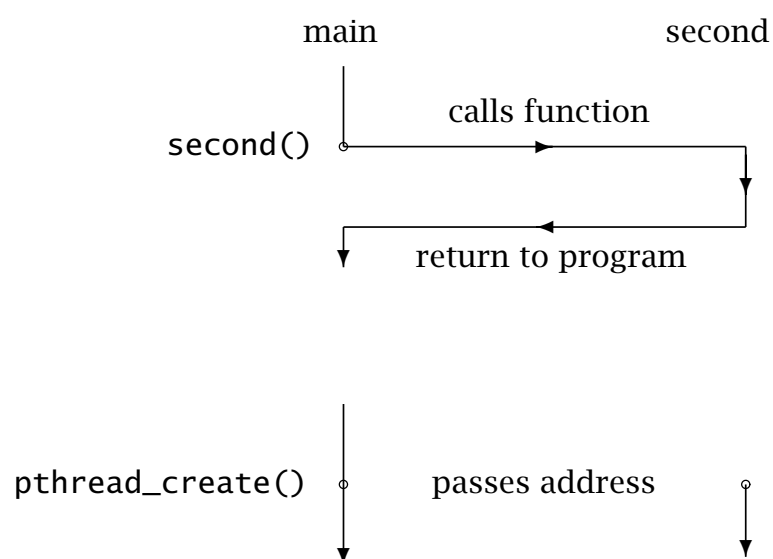
Figure 3.6: How the system keeps track of processes

```
struct task_struct{
    volatile long      state;
    long               counter, priority;
    struct task_struct *next_task, *prev_task;
    struct task_struct *next_run, *prev_run;
    int                exit_code, exit_signal;
    int                pid;
    struct task_struct *p_opptr, *p_cptra;
    struct wait_queue  *wait_chldexit;
    struct task_struct *pidhash_next;
    unsigned long      policy;
    struct tms          times;
    unsigned long      start_time;
    unsigned short     uid, gid;
    struct thread_struct tss;
    struct files_struct *files;
    struct mm_struct    *mm;
    struct signal_struct *sig;
    sigset_t            signal, blocked;
};
```

**Figure 3.7:** Data structure representing a process

```
printf ("Before the fork\n");  
fork ();  
printf ("After the fork\n");
```

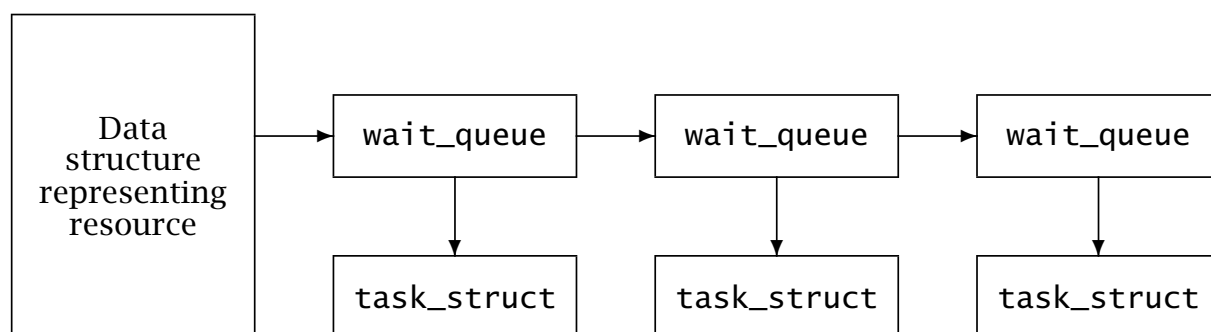
**Figure 3.8:** Program to illustrate fork()



**Figure 3.9:** Function call and thread creation

```
struct wait_queue{  
    struct task_struct *task;  
    struct wait_queue *next;  
};
```

**Figure 3.10:** An entry in a wait queue



**Figure 3.11:** Three processes waiting for a resource

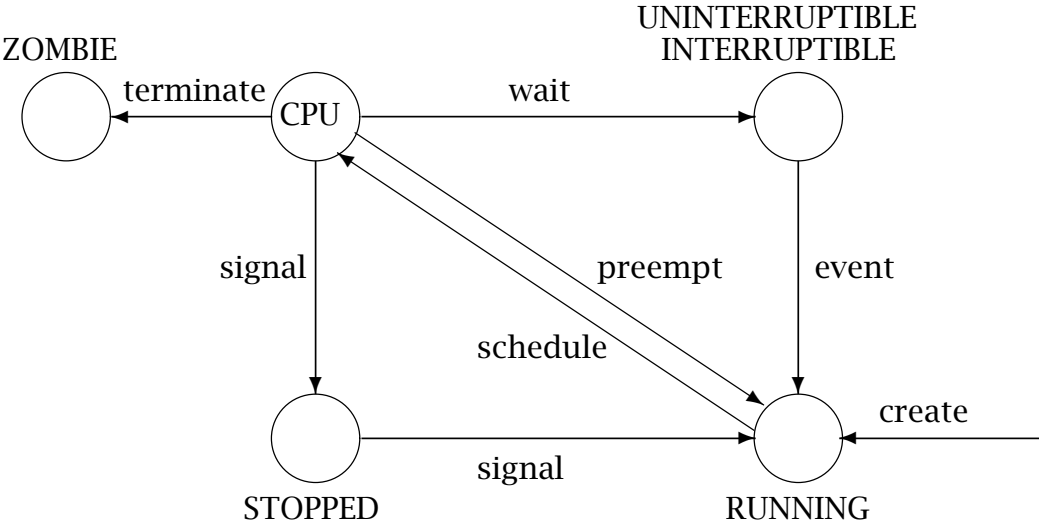


Figure 3.12: Process states and transitions



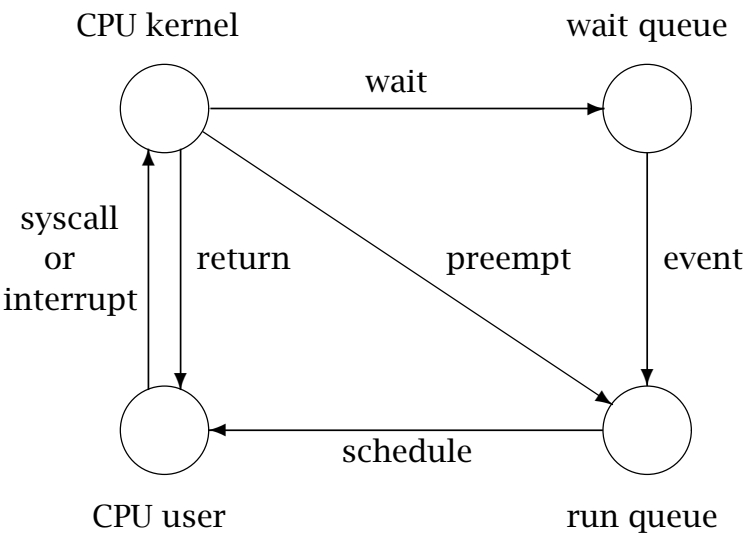


Figure 3.13: Running in user or kernel mode

```
struct thread_struct{
    unsigned long  cr3;
    unsigned long  eip;
    unsigned long  eflags;
    unsigned long  eax,ebx,ecx,edx;
    unsigned long  esp;
    unsigned long  ebp;
    unsigned long  esi;
    unsigned long  edi;
    unsigned short ss;
    unsigned short cs;
    unsigned short ds;
    unsigned short es;
    unsigned short fs;
    unsigned short gs;
};
```

**Figure 3.14:** The volatile environment of a process


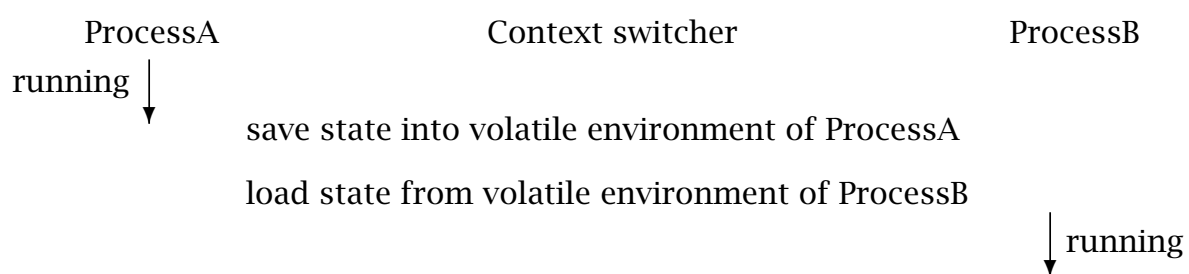



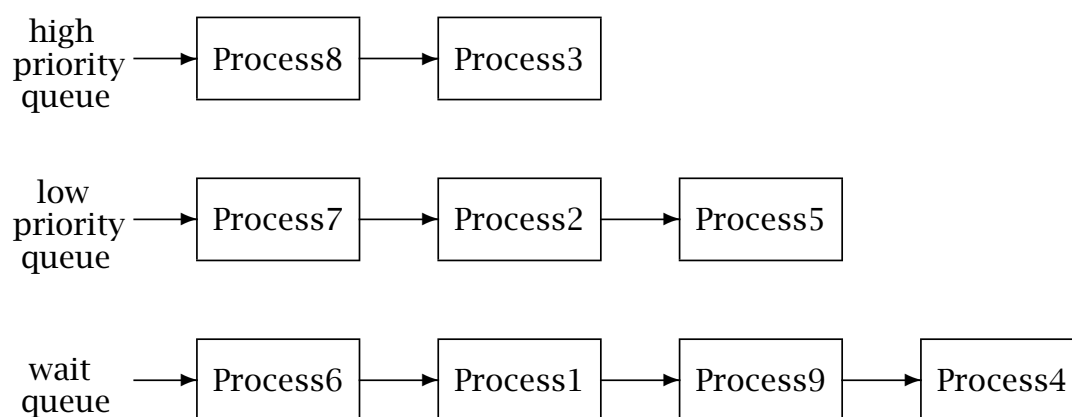
Figure 3.15: Data saved by context switcher



**Figure 3.16:** Context switching



**Figure 3.17:** Processes on the ready and wait queues



**Figure 3.18:** Two different priority queues

```
MOV EAX, items ; load from items into register EAX
INC EAX        ; increment register EAX
MOV items, EAX ; store from register EAX back to items
```

**Figure 4.1:** Code to increment counter

```
MOV EAX, items ; load from items to register EAX
DEC EAX        ; decrement register EAX
MOV items, EAX ; store from register EAX back to items
```

**Figure 4.2:** Code to decrement counter



beginning section  
Entry protocol for critical section  
critical section  
Exit protocol for critical section  
remainder section

**Figure 4.3:** General structure of a cooperating program

```
beginning section

WHILE (guard == 1)
    DO nothing
ENDWHILE
guard = 1

critical section

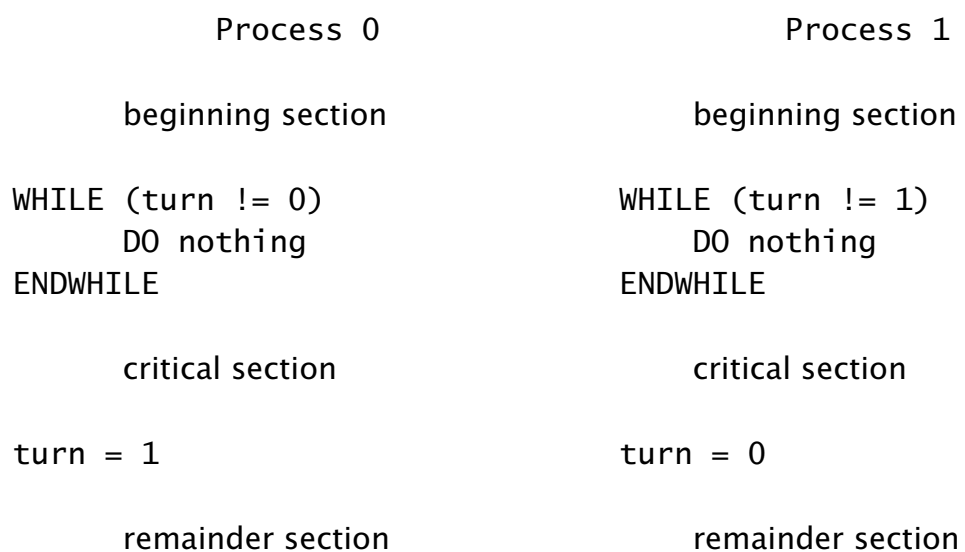
guard = 0

remainder section
```

**Figure 4.4:** First attempt at solution for two processes

```
test: MOV EAX, guard
      CMP EAX, #1
      JE  test:
      MOV guard, #1
```

**Figure 4.5:** Compiled version of entry protocol



**Figure 4.6:** Algorithms using turn

```
beginning section  
  
WHILE (flag[1] == 1)  
    DO nothing  
ENDWHILE  
flag[0] = 1  
  
critical section  
  
flag[0] = 0  
  
remainder section
```

**Figure 4.7:** Algorithm using two flags

```
beginning section

flag[0] = 1
WHILE (flag[1] == 1)
    DO nothing
ENDWHILE

critical section

flag[0] = 0

remainder section
```

**Figure 4.8:** Setting the flag before testing

```
beginning section

flag[0] = 1
turn = 1
WHILE ((flag[1] == 1) AND (turn == 1))
    DO nothing
ENDWHILE

critical section

flag[0] = 0


remainder section
```

**Figure 4.9:** Algorithm for mutual exclusion

Process 0	Process 1
MOV flag[0], #1	MOV flag[1], #1
MOV turn, #1	MOV turn, #0
test: CMP flag[1], #1	test: CMP flag[0], #1
JNE enter:	JNE enter:
CMP turn, #1	CMP turn, #0
JE test:	JE test:
enter:	enter:

**Figure 4.10:** Entry protocols for both processes





**Figure 4.11:** Eight contending processes

```
beginning section

REPEAT
    flag[4] = want-in
    p = turn
    WHILE (p ≠ 4)
        IF (flag[p] == idle) THEN
            p++ MOD 8
        ELSE
            p = turn
        ENDIF
    ENDWHILE
    flag[4] = in-cs

    j = 0
    WHILE ((j < 8) AND ((j == 4) OR (flag[j] ≠ in-cs)))
        j ++
    ENDWHILE
UNTIL ((j == 8) AND ((turn == 4) OR (flag[turn] == idle)))
turn = 4

critical section

flag[4] = idle

remainder section
```

**Figure 4.12:** Eisenberg and McGuire algorithm

```
beginning section

choosing[i] = TRUE
counter++
number[i] = counter
choosing[i] = FALSE

FOR j = 0 TO n - 1 DO
    WHILE (choosing[j] == TRUE)
        DO nothing
    ENDWHILE
    WHILE ((number[j] ≠ 0) AND (number[j] < number[i]))
        DO nothing
    ENDWHILE
ENDFOR

critical section

number[i] = 0

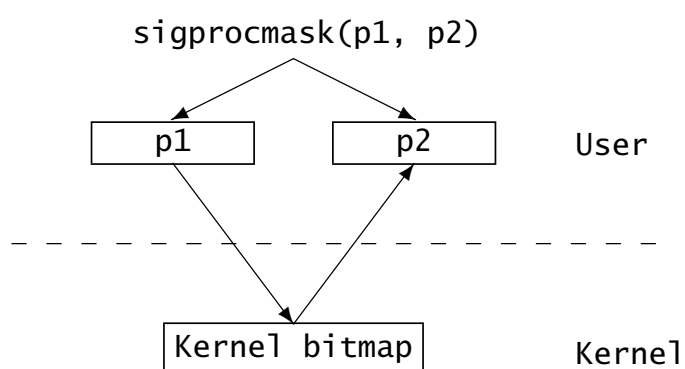
remainder section
```

**Figure 4.13:** Lamport’s algorithm



```
#define SIGINT    2  /* interrupt, generated from terminal */
#define SIGILL    4  /* illegal instruction */
#define SIGABRT   6  /* abort process */
#define SIGFPE    8  /* floating point exception */
#define SIGKILL   9  /* kill a process */
#define SIGUSR1  10  /* user defined signal 1 */
#define SIGSEGV  11  /* segmentation violation */
#define SIGUSR2  12  /* user defined signal 2 */
#define SIGALRM  14  /* alarm clock timeout */
#define SIGCHLD  17  /* sent to parent on child exit */
#define SIGXCPU  24  /* cpu time limit exceeded */
```

**Figure 5.1:** A selection of signal values



**Figure 5.2:** Changing the signal mask

```
struct signal_struct{
    atomic_t          count;
    struct k_sigaction action[32];
    spinlock_t        siglock;
};
```

**Figure 5.3:** Data structure tracking signal handlers

```
WHILE (Test_and_set (Key))  
    DO nothing  
ENDWHILE
```

**Figure 5.4:** Entry protocol using test and set



```
REPEAT
    Local = 1
    Exchange(Local, Key)
UNTIL Local == 0
```

**Figure 5.5:** Entry protocol using exchange

beginning section

WAIT(Guard)

critical section

SIGNAL(Guard)

remainder section

**Figure 5.6:** Semaphore for mutual exclusion




Figure 5.7: Successive states of a mutual exclusion semaphore




Figure 5.8: Semaphore for synchronisation




Figure 5.9: An alternative for synchronisation




Figure 5.10: Semaphore for resource allocation




Figure 5.11: Circular buffer with eight slots

```
Produce an item
WAIT(SlotFree)
Put item in buffer at NextIn
NextIn++
SIGNAL(ItemAvailable)
```

**Figure 5.12:** Algorithm for producer



```
WAIT(ItemAvailable)
Get item from buffer at NextOut
NextOut++
SIGNAL(SlotFree)
Consume the item
```

**Figure 5.13:** Algorithm for consumer

```
Produce an item
WAIT(SlotFree)
WAIT(Guard)
Put item in buffer at NextIn
NextIn++
SIGNAL(Guard)
SIGNAL(ItemAvailable)
```

**Figure 5.14:** Algorithm for one of many producers

```
WAIT(ItemAvailable)
WAIT(Guard)
Get item from buffer at NextOut
NextOut++
SIGNAL (Guard)
SIGNAL(SlotFree)
Consume the item
```

**Figure 5.15:** Algorithm for one of many consumers

```
WAIT(Guard)
readers++
IF (readers == 1) THEN
    WAIT(Writing)
ENDIF
SIGNAL(Guard)
```

**Figure 5.16:** Entry protocol for a reader

```
WAIT(Guard)
readers--
IF (readers == 0) THEN
    SIGNAL(Writing)
ENDIF
SIGNAL(Guard)
```

**Figure 5.17:** Exit protocol for a reader

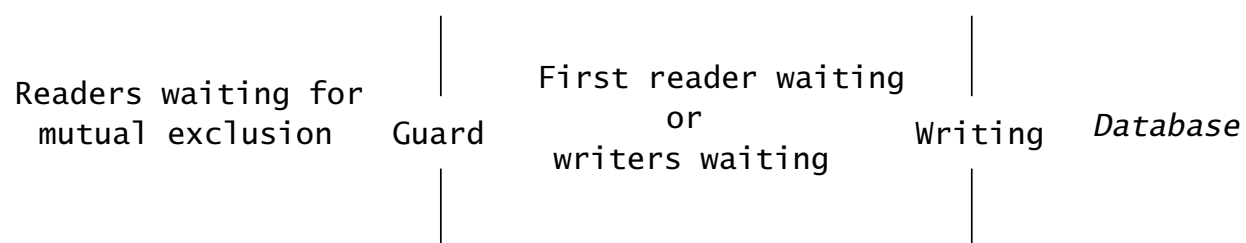


Figure 5.18: Solution 2




Figure 5.19: Solution 3

```
WAIT(Guard)
wr++
IF (ww == 0) THEN
    ar++
    SIGNAL(R)
ENDIF
SIGNAL(Guard)
WAIT (R)
```

**Figure 5.20:** Reader prologue code



```
WAIT(Guard)
wr++
IF (ww > 0) THEN
    WAIT(R)
ELSE
    ar++
ENDIF
SIGNAL(Guard)
```

**Figure 5.21:** Incorrect reader prologue

```
WAIT (Guard)
ar--
wr--
IF (ar == 0) THEN
    IF (ww > 0) THEN
        SIGNAL(w)
    ENDIF
ENDIF
SIGNAL (Guard)
```

**Figure 5.22:** Reader epilogue code

```
WAIT(Guard)
ww++
IF (ar == 0) AND (ww == 1) THEN
    SIGNAL(W)
ENDIF
SIGNAL(Guard)
WAIT(W)
```

**Figure 5.23:** Writer prologue code

```
WAIT(Guard)
ww--
WHILE (ar < wr)
    ar++
    SIGNAL(R)
ENDWHILE
IF (ar == 0) AND (ww > 0) THEN
    SIGNAL(W)
ENDIF
SIGNAL(Guard)
```

**Figure 5.24:** Writer epilogue code

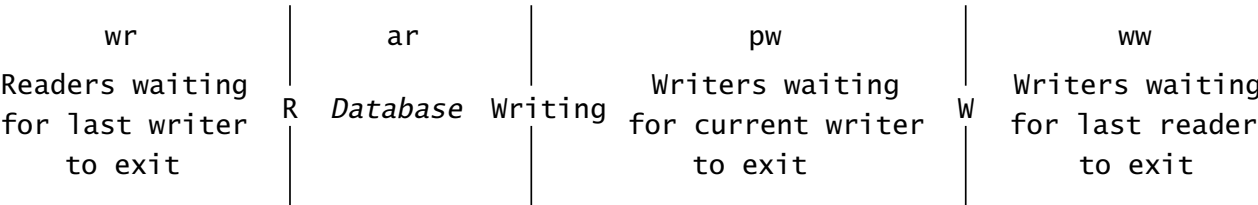


Figure 5.25: Solution 4

```
WAIT (Guard)
ar--
wr--
IF (ar == 0) THEN
    WHILE (pw < ww)
        pw++
        SIGNAL(W)
    ENDWHILE
ENDIF
SIGNAL (Guard)
```

**Figure 5.26:** Reader epilogue code

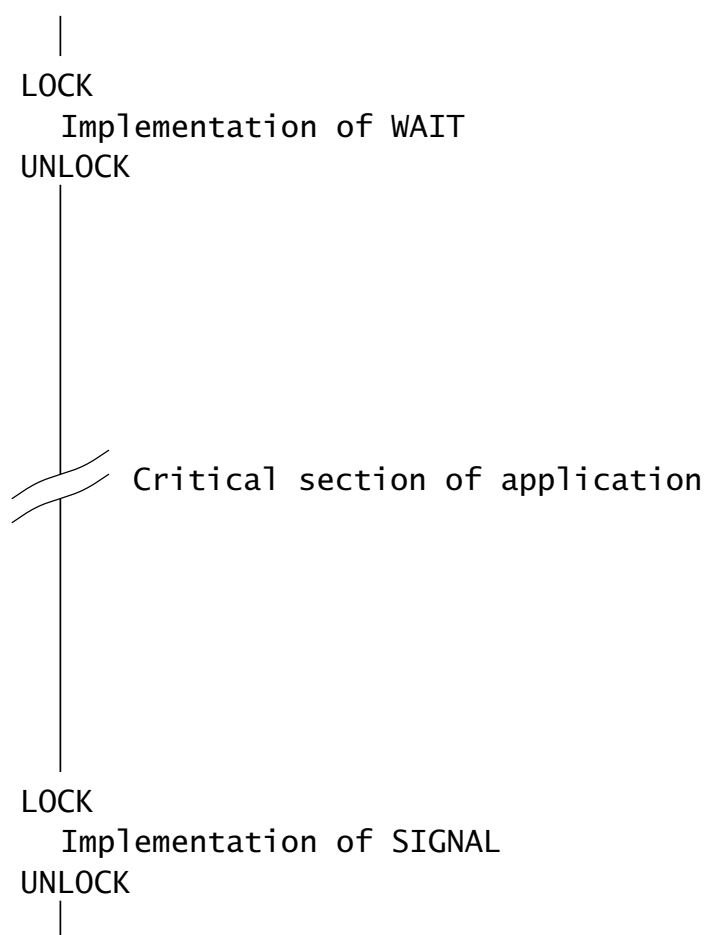
```
WAIT(Guard)
ww++
IF (ar == 0) THEN
    pw++
    SIGNAL(W)
ENDIF
SIGNAL(Guard)
WAIT(W)
WAIT(Writing)
```

**Figure 5.27:** Writer prologue code

```
SIGNAL(Writing)
WAIT(Guard)
pw--
ww--
IF (pw == 0) THEN
    WHILE (ar < wr)
        ar++
        SIGNAL(R)
    ENDWHILE
ENDIF
SIGNAL(Guard)
```

**Figure 5.28:** Writer epilogue code





**Figure 5.29:** Timing relationship of locks and semaphores

```
struct semid_ds{  
    struct ipc_perm  sem_perm;  
    struct sem        *sem_base;  
    struct sem_queue  *sem_pending;  
};
```

**Figure 5.30:** Data structure representing a semaphore set

```
struct ipc_perm{
    kernel_key_t  key; /* user supplied key */
    kernel_uid_t  uid; /* owner's user id   */
    kernel_gid_t  gid; /* owner's group id  */
    kernel_mode_t mode; /* access modes      */
};
```

**Figure 5.31:** Data structure controlling access to a semaphore

```
struct sem{  
    int semval; /* current value */  
    int sempid; /* process which last operated on sem */  
};
```

**Figure 5.32:** Data structure representing an individual semaphore


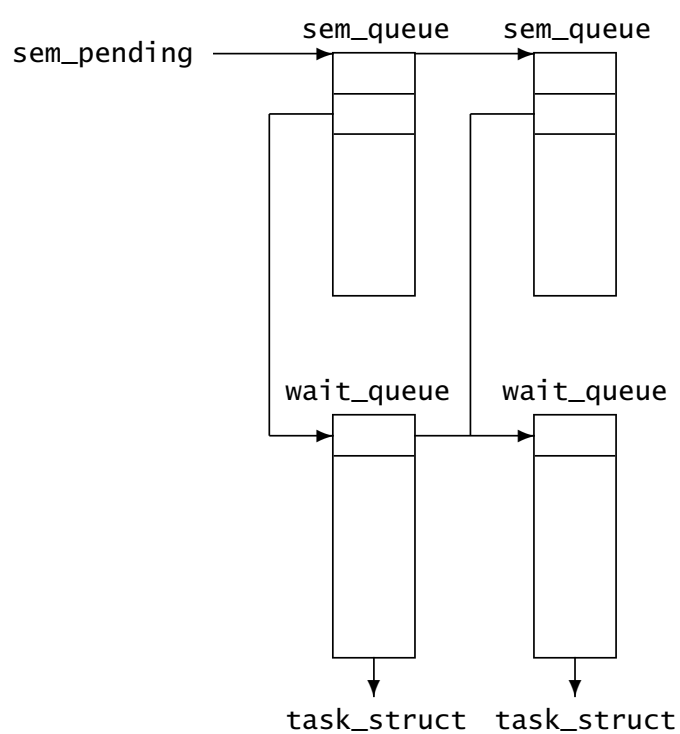


Figure 5.33: System V semaphore layout

```
struct sem_queue{
    struct sem_queue *next;
    struct wait_queue *sleeper;
    struct semid_ds *sma;
    struct sembuf *sops;
    int nsops;
};
```

**Figure 5.34:** Data structure representing a waiting process



**Figure 5.35:** Processes waiting on the semaphore set

Producer	Consumer
ADVANCE(E)	AWAIT(E, i)
	i++

**Figure 5.36:** Eventcount for synchronisation




beginning section  
AWAIT(E, TICKET(S))  
critical section  
ADVANCE(E)  
remainder section


**Figure 5.37:** Eventcount and sequencer for mutual exclusion

Producer	Consumer
AWAIT(OUT, $i + 1 - N$ )	AWAIT(IN, $j + 1$ )
insert at $i \text{ MOD } N$	Remove from $j \text{ MOD } N$
$i++$	$j++$
ADVANCE(IN)	ADVANCE(OUT)

**Figure 5.38:** Eventcounts with one producer, one consumer



**Figure 6.1:** Outline of message passing



**Figure 6.2:** A message queue

```
msqid_ds{
    struct ipc_perm    msg_perm;    /* access permissions        */
    struct msg         *msg_first;  /* first message on queue    */
    struct msg         *msg_last;   /* last message on queue     */
    struct wait_queue  *wwait;      /* blocked writing threads    */
    struct wait_queue  *rwait;      /* blocked reading threads   */
    unsigned short     msg_qnum;    /* number of messages on queue */
};
```

**Figure 6.3:** Data structure representing a message queue

```
struct msg{
    struct msg *msg_next; /* next message on queue */
    long      msg_type; /* as specified by sender */
    char      *msg_spot; /* message text address */
    time_t    msg_stime; /* msgsnd time */
    short     msg_ts; /* message text size */
};
```

**Figure 6.4:** Structure of an actual message

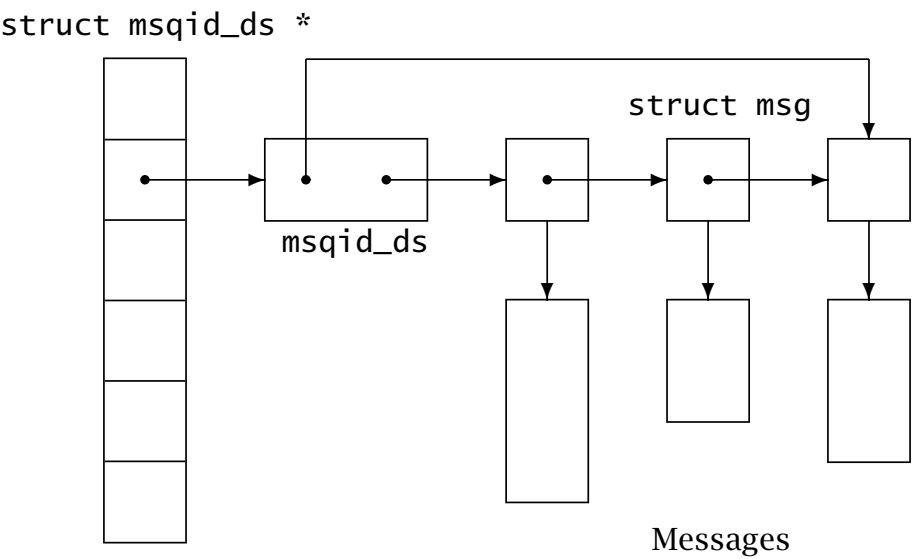


Figure 6.5: System V message queue layout

```
struct msgbuf{  
    long mtype; /* message type */  
    char mtext[]; /* message contents */  
};
```

**Figure 6.6:** Data structure representing a message



```
shared struct{
    int writers;
    int readers;
}shared_counts;

shared char writelock;

/* reader code */

critical region shared_counts{
    await (writers == 0);
    readers++;
}

    /* READ */

critical region shared_counts{
    readers--;
}

/* writer code */

critical region shared_counts{
    writers++;
    await (readers == 0);
}

critical region writelock{

    /* WRITE */

}

critical region shared_counts{
    writers--;
}
```

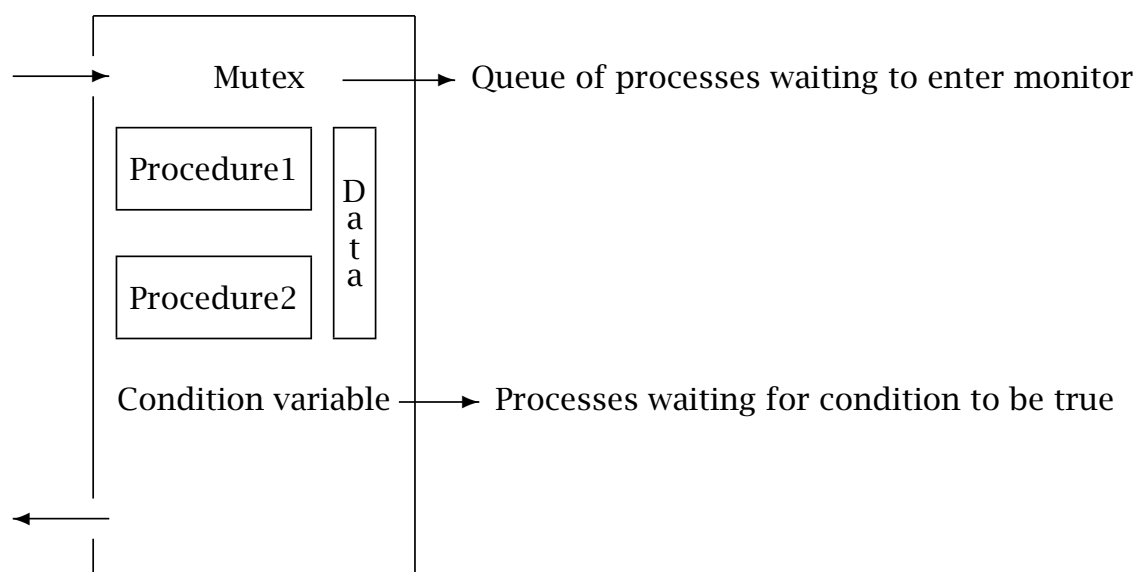
**Figure 6.7:** Readers/writers using conditional critical regions

```
CWAIT(conditionvar)
    LOCK
    Mark process unrunnable
    Queue the process on the condition variable
    Release mutual exclusion on the monitor
    UNLOCK
    Call context switcher
```

**Figure 6.8:** Implementation of CWAIT

```
CSIGNAL(conditionvar)
  LOCK
  IF there is a process waiting on conditionvar THEN
    Mark one runnable
    Mark current process unrunnable
    Transfer mutual exclusion on the monitor to selected process
  UNLOCK
  Call context switcher
ELSE
  UNLOCK
  Return
ENDIF
```

**Figure 6.9:** Implementation of CSIGNAL



**Figure 6.10:** An illustration of a monitor

```
MONITOR allocator{

    boolean busy = FALSE;
    condition free;

    reserve(){
        while (busy == TRUE)
            CWAIT(free);
        busy = TRUE;
    }

    release(){
        busy = FALSE;
        CSIGNAL(free);
    }
}
```

**Figure 6.11:** Monitor to allocate a single resource

```
MONITOR  buffer{

    int count = 0;
    condition spaceavail;
    condition itemavail;

    producer(){
        while (count == MAX)
            CWAIT(spaceavail);
        /* Put item in buffer */
        count++;
        CSIGNAL(itemavail);
    }

    consumer(){
        while (count == 0)
            CWAIT(itemavail);
        /* Get item from buffer */
        count--;
        CSIGNAL(spaceavail);
    }
}
```

**Figure 6.12:** Monitor to manage a bounded buffer

```
MONITOR  readers-writers{

    int      readers = 0, writers = 0;
    boolean  busy-writing = FALSE;
    condition readers-waiting, writers-waiting;

    StartRead(){
        while (writers > 0)
            CWAIT(writers-waiting);
        readers++;
        CSIGNAL(readers-waiting)
    }

    EndRead(){
        readers--;
        if (readers == 0)
            CSIGNAL(writers-waiting)
    }


    StartWrite(){
        writers++;
        while ((busy-writing == TRUE) || (readers > 0))
            CWAIT(writers-waiting);
        busy-writing = TRUE
    }

    EndWrite(){
        busy-writing = FALSE;
        writers--;
        if (writers > 0)
            CSIGNAL(writers-waiting);
        else
            CSIGNAL(readers-waiting);
    }
}
```


**Figure 6.13:** Monitor to implement reader/writer interlock








**Figure 7.1:** A resource allocation graph



**Figure 7.2:** A resource allocation graph for four processes



**Figure 7.3:** Safe, unsafe, and deadlocked system




Figure 7.4: A claims graph




Figure 7.5: One resource allocated

```
IF (Request > Need[i]) THEN
    Error-illegal request
ENDIF
IF (Request > Available) THEN
    Wait
ENDIF

Available = Available - Request
Allocation[i] = Allocation[i] + Request
Need[i] = Need[i] - Request

Check if this is a safe state
IF Safe THEN
    Allocate resources
ELSE
    Restore state
    Wait
ENDIF
```

**Figure 7.6:** The banker’s algorithm

```
Work[r] = Available[r]
Finish[p] = FALSE (all elements)

REPEAT
    Found = FALSE
    i = 0
    REPEAT
        IF (Finish[i] == FALSE) AND (Need[i] ≤ Work) THEN
            Finish[i] = TRUE
            Work = Work + Allocation[i]
            Found = TRUE
        ENDIF
        i++
    UNTIL (Found == TRUE) OR (i == p)
UNTIL Found == FALSE

Safe = TRUE
FOR i = 0 TO p - 1
    IF (Finish[i] == FALSE) THEN
        Safe = FALSE
    ENDIF
ENDFOR
```

**Figure 7.7:** Safety algorithm

Available						
		A	B			
		3	5			

	Max		Allocation		Need	
	A	B	A	B	A	B
$P_0$	7	5	0	1	7	4
$P_1$	3	2	2	0	1	2
$P_2$	9	0	3	0	6	0
$P_3$	2	2	2	1	0	1

**Figure 7.8:** State of the system before request



OUTER LOOP		INNER LOOP	
1	Work = 3,5	i = 0    Need[0] > Work i = 1    Need[1] < Work	Finish[1] = TRUE
2	Work = 5,5	i = 0    Need[0] > Work i = 1    Finish[1] = TRUE i = 2    Need[2] > Work i = 3    Need[3] < Work	Finish[3] = TRUE
3	Work = 7,6	i = 0    Need[0] < Work	Finish[0] = TRUE
4	Work = 7,7	i = 0    Finish[0] = TRUE i = 1    Finish[1] = TRUE i = 2    Need[2] < Work	Finish[2] = TRUE
5	Work = 10,7	i = 0    Finish[0] = TRUE i = 1    Finish[1] = TRUE i = 2    Finish[2] = TRUE i = 3    Finish[3] = TRUE	

**Figure 7.9:** Application of the safety algorithm

Available							
		A		B			
		<u>2</u>		5			
		Max		Allocation		Need	
		A	B	A	B	A	B
$P_0$		7	5	0	1	7	4
$P_1$		3	2	<u>3</u>	0	<u>0</u>	2
$P_2$		9	0	3	0	6	0
$P_3$		2	2	2	1	0	1

Figure 7.10: State of the system reflecting the request




Figure 7.11: Resource allocation graph




Figure 7.12: Wait-for graph

```
FOR i = 0 TO MaxProc - 1
    IF (Allocation[i] == 0) THEN
        Finish[i] = TRUE
    ELSE
        Finish[i] = FALSE
    ENDIF
ENDFOR

REPEAT
    Found = FALSE
    i = 0
    REPEAT
        IF (Finish[i] == FALSE) AND (Request[i] ≤ Work) THEN
            Work = Work + Allocation[i]
            Finish[i] = TRUE
            Found = TRUE
        ENDIF
        i++
    UNTIL (Found == TRUE) OR (i == p)
UNTIL (Found == FALSE)

Deadlocked = FALSE
FOR i = 0 TO MaxProc - 1
    IF (Finish[i] == FALSE) THEN
        Deadlocked = TRUE
    ENDIF
ENDFOR
```

**Figure 7.13:** Algorithm to detect deadlock

Available					
		A	B		
		0	0		

	Allocation		Request	
	A	B	A	B
$P_0$	0	1	0	0
$P_1$	2	0	2	0
$P_2$	3	0	0	0
$P_3$	2	1	1	0

Figure 7.14: A system which is not deadlocked

OUTER LOOP			INNER LOOP	
1	Work = 0,0	i = 0	Request[0] < Work	Finish[0] = TRUE
2	Work = 0,1	i = 0	Finish[0] = TRUE	
		i = 1	Request[1] > Work	
		i = 2	Request[2] < Work	Finish[2] = TRUE
3	Work = 3,1	i = 0	Finish[0] = TRUE	
		i = 1	Request[1] < Work	Finish[1] = TRUE
4	Work = 5,1	i = 0	Finish[0] = TRUE	
		i = 1	Finish[1] = TRUE	
		i = 2	Finish[2] = TRUE	
		i = 3	Request[3] < Work	Finish[3] = TRUE
FINAL LOOP		i = 0	Finish[0] = TRUE	
		i = 1	Finish[1] = TRUE	
		i = 2	Finish[2] = TRUE	
		i = 4	Finish[3] = TRUE	Deadlocked = FALSE

**Figure 7.15:** Application of the deadlock detection algorithm

Available					
		A	B		
		0	0		

	Allocation		Request	
	A	B	A	B
$P_0$	0	1	0	0
$P_1$	2	0	2	0
$P_2$	3	0	<u>1</u>	0
$P_3$	2	1	1	0

Figure 7.16: A deadlocked system






Figure 8.1: Hardware memory

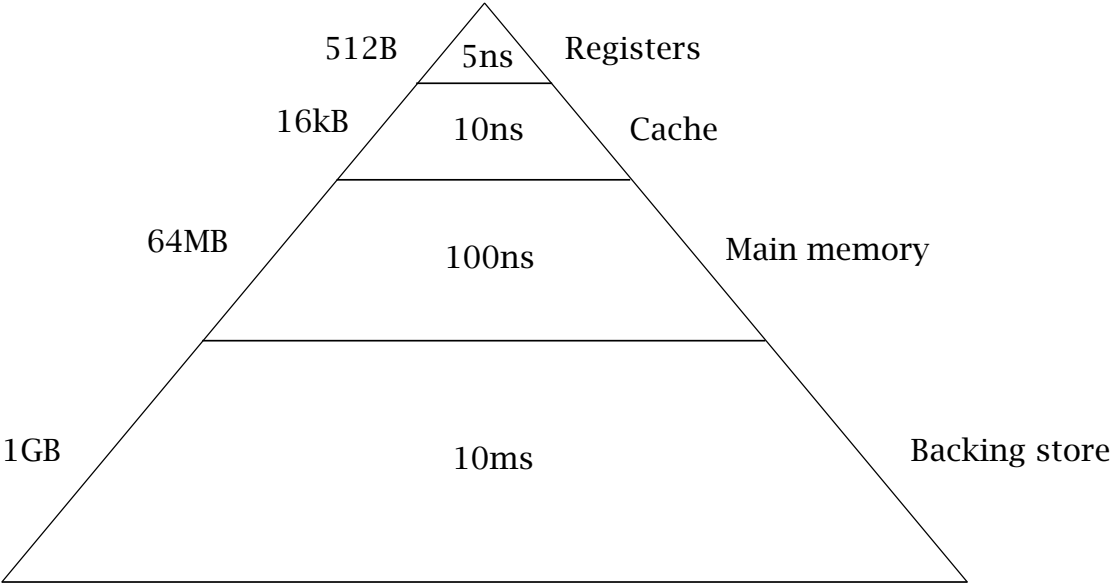


Figure 8.2: The memory pyramid

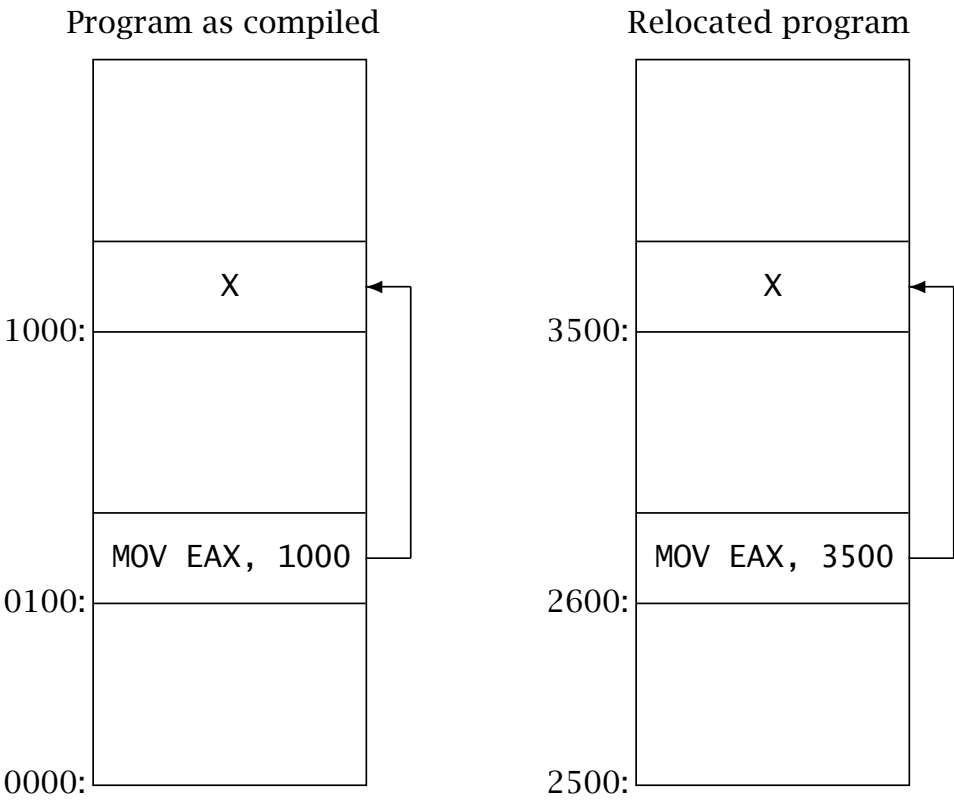




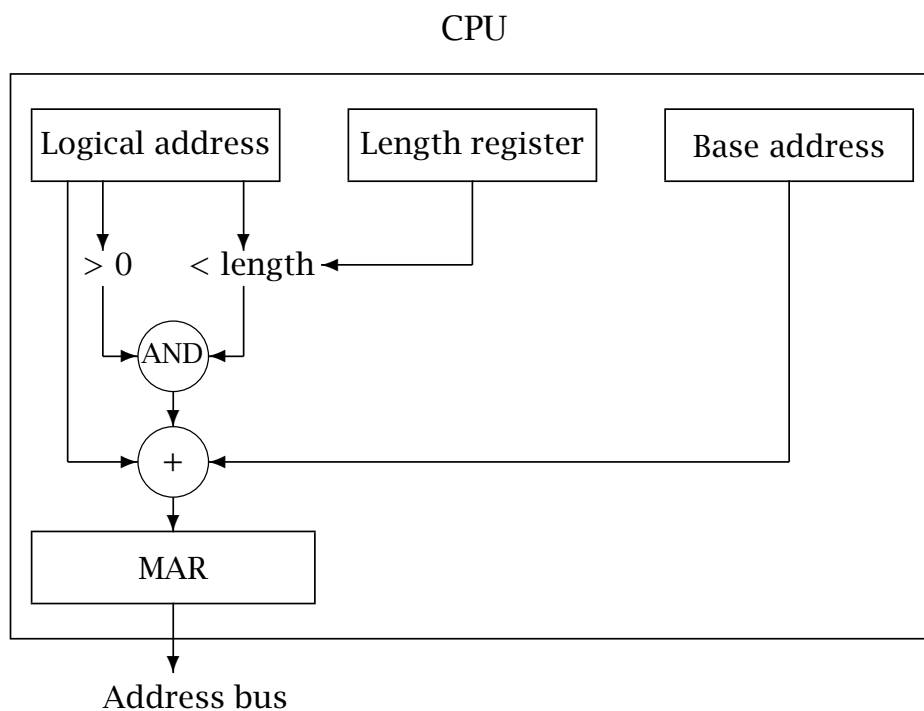
Figure 8.3: Program as loaded by relocating loader



**Figure 8.4:** Virtual memory



**Figure 8.5:** Address modification in CPU



**Figure 8.6:** Protection and address modification

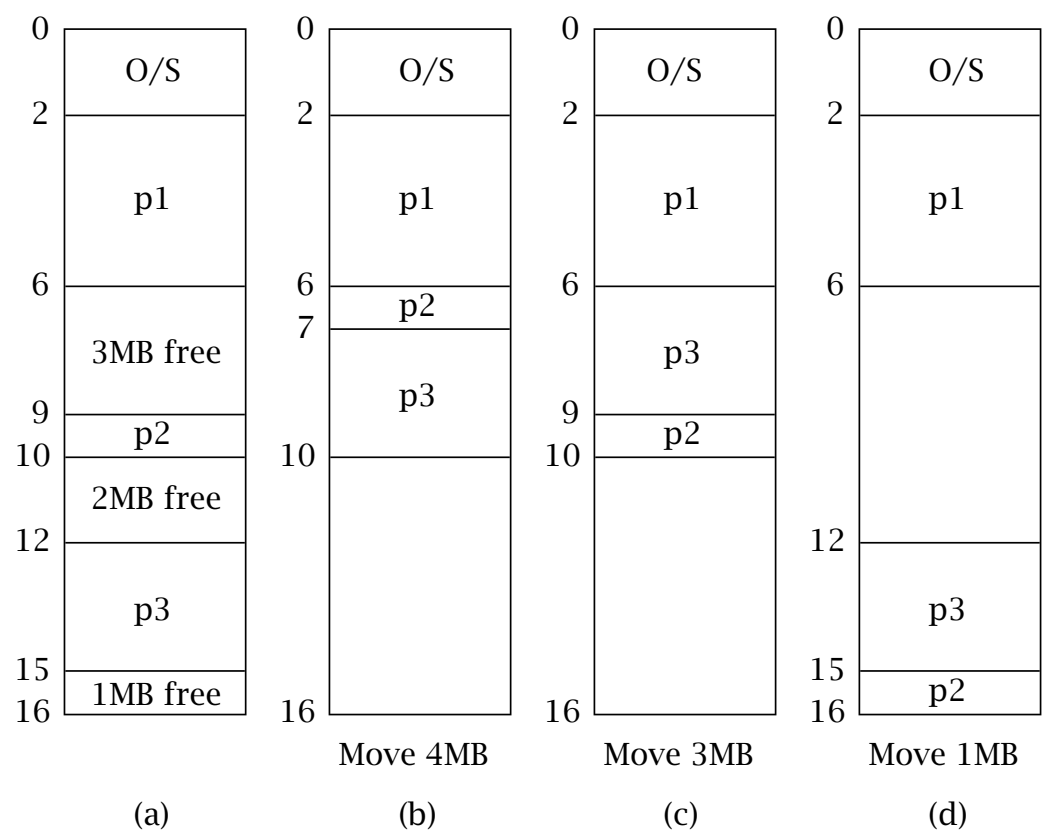



Figure 8.7: Compaction



**Figure 8.8:** A logical address



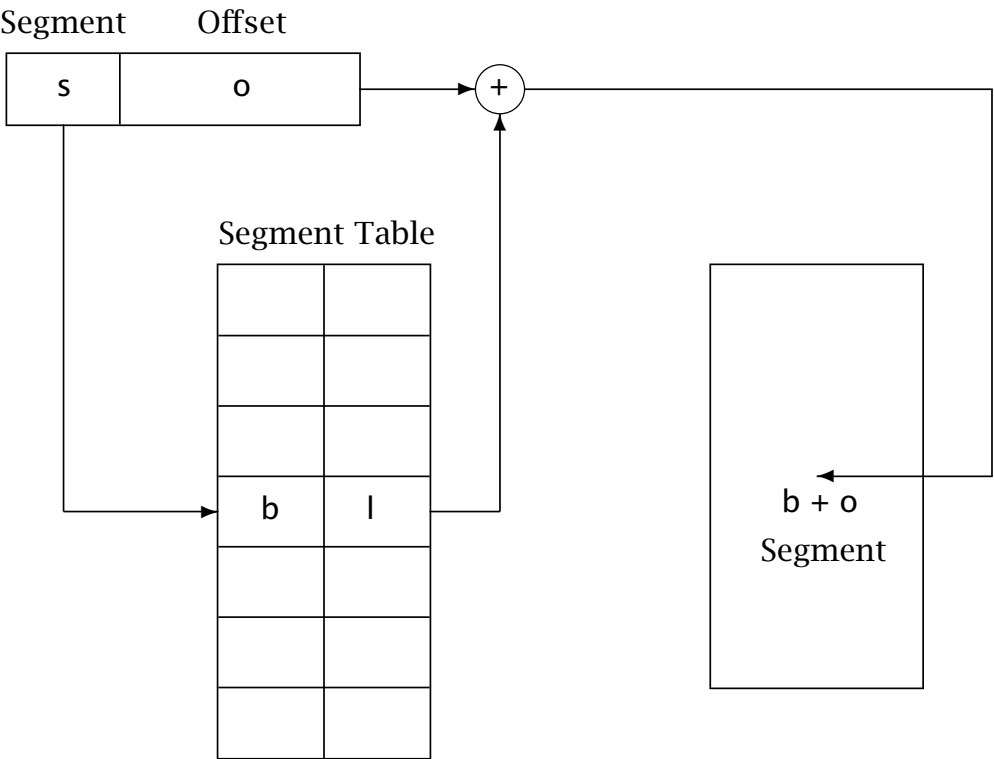


Figure 8.9: Address translation with segment table

r	w	e	base address	length
---	---	---	--------------	--------

**Figure 8.10:** A segment table entry

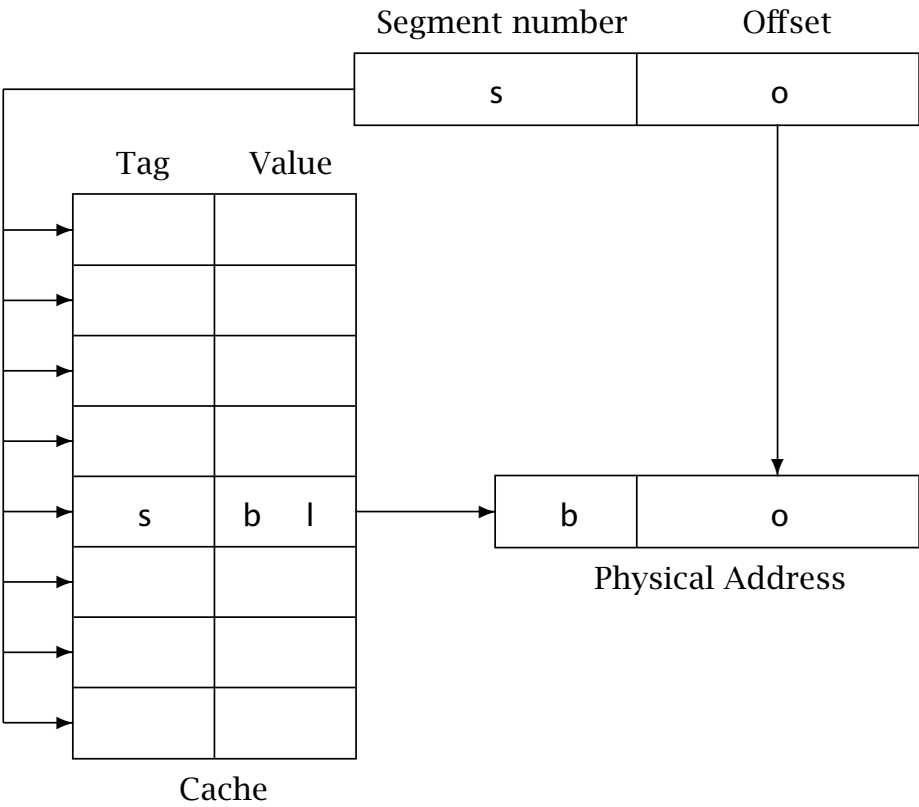


Figure 8.11: Searching a segmentation cache

```
Present segment number to associative store
IF no match, THEN
    Use segment number as index into segment table
    IF not present, THEN
        Segment fault—fetch from secondary memory
        Update segment table
    ENDIF
    Update associative store
ENDIF
Add offset to base address
```

**Figure 8.12:** Segmentation algorithm

```
struct page{
    struct page    *next;
    struct inode    *inode;
    unsigned long  offset;
    struct page    *next_hash;
    atomic_t       count;
};
```

**Figure 8.13:** Data structure representing a physical page

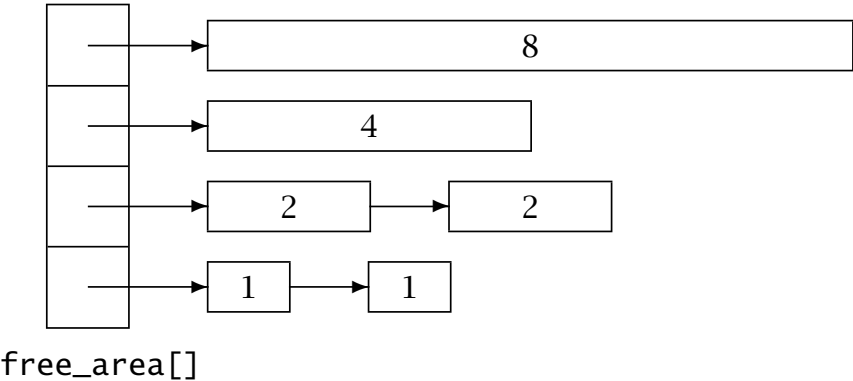


Figure 8.14: Tracking free physical memory


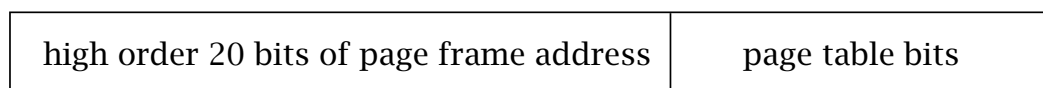


Figure 8.15: Linux page table structure



**Figure 8.16:** Format of page table entry



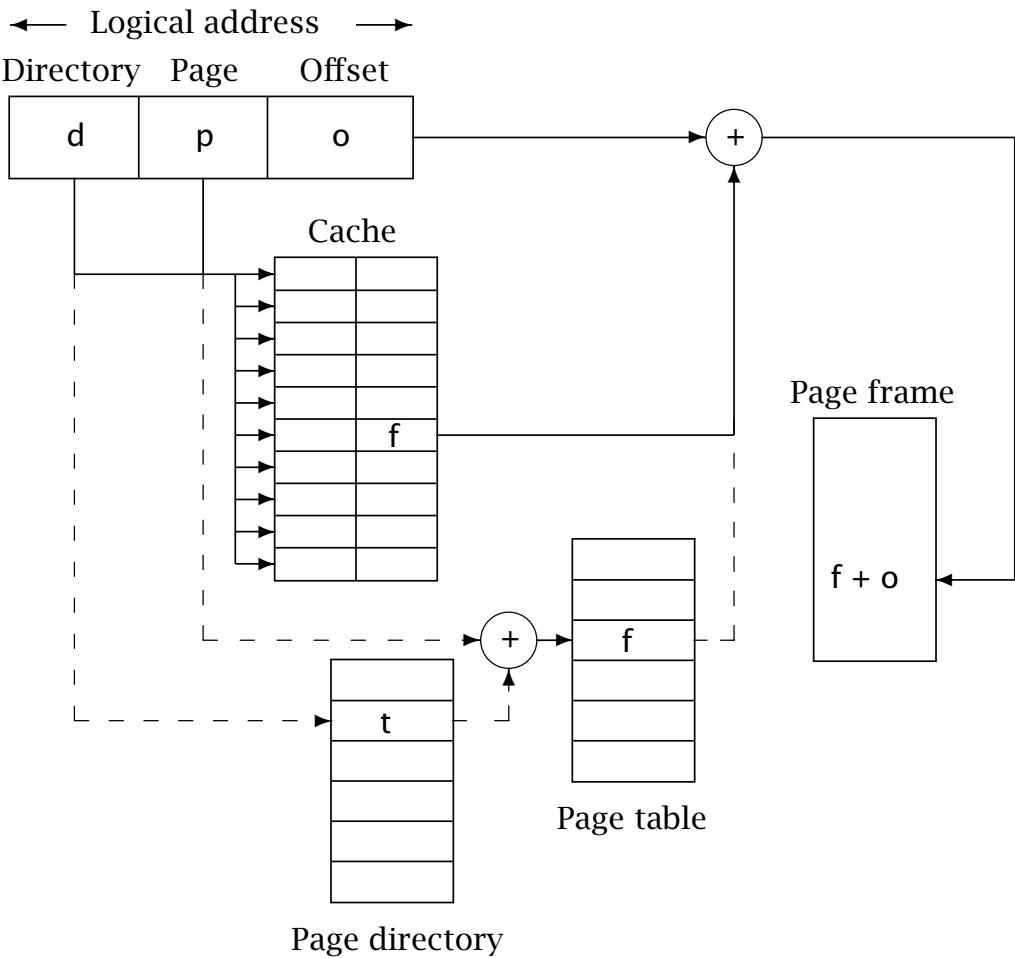


Figure 8.17: Address translation with cache and page table

```
int  array[128][128];  
  
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        array[i][j] = 0;
```

**Figure 8.18:** Program to initialise an array

```
struct mm_struct{
    struct vm_area_struct *mmap;
    pgd_t                *pgd;
    int                  count;
    unsigned long         start_code, end_code;
    unsigned long         start_data, end_data;
    unsigned long         start_stack;
};
```

**Figure 8.19:** The root of the memory management data structures

```
struct vm_area_struct{
    struct mm_struct      *vm_mm;
    unsigned long         vm_start;
    unsigned long         vm_end;
    struct vm_area_struct *vm_next;
    unsigned short         vm_flags;
    struct vm_operations_struct *vm_ops;
    unsigned long         vm_offset;
    struct file            *vm_file;
};
```

**Figure 8.20:** Data structure controlling a region




Figure 8.21: Memory management data structures for a process

```
#define VM_READ      0x0001
#define VM_WRITE     0x0002
#define VM_EXEC      0x0004
#define VM_SHARED     0x0008
#define VM_LOCKED    0x2000
```

**Figure 8.22:** Possible values for the flags field

```
struct vm_operations_struct{  
    void      (*open)();  
    void      (*close)();  
    void      (*unmap)();  
    int       (*swapout)();  
    pte_t     (*swapin)();  
};
```

**Figure 8.23:** Operations on a region




Figure 8.24: Pages actually in memory



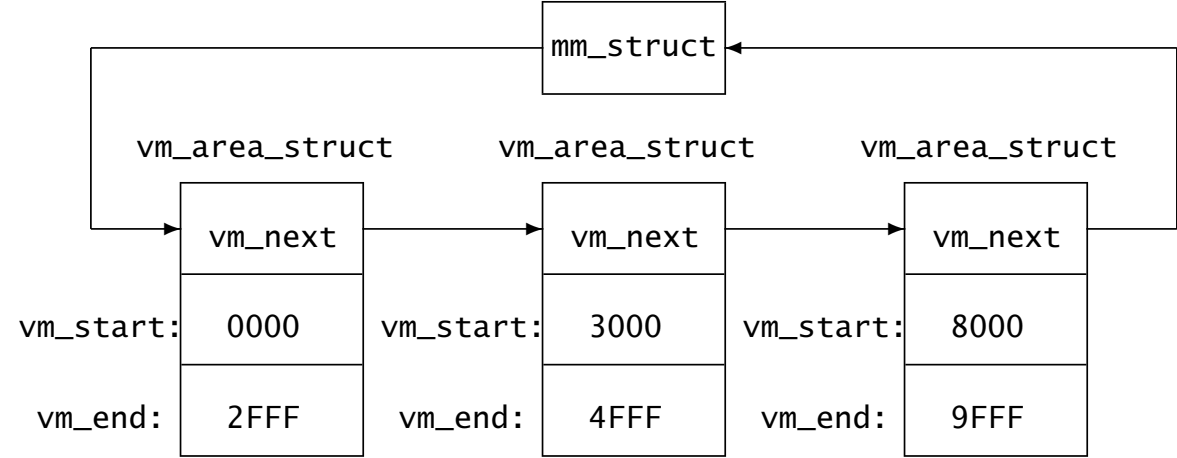




Figure 8.25: Example mapping of three regions




**Figure 8.26:** Relationship between page frames and page faults

```
REPEAT
  Examine the reference bit in the page table entry for pagenumber
  IF (reference bit == 1) THEN
    reference bit = 0
  ELSE
    remove page
  ENDIF
  pagenumber = (pagenumber + 1) MOD size of page table
UNTIL (required number of pages removed)
```


**Figure 8.27:** Algorithm for NRU replacement



**Figure 8.28:** Effect of thrashing on CPU utilisation



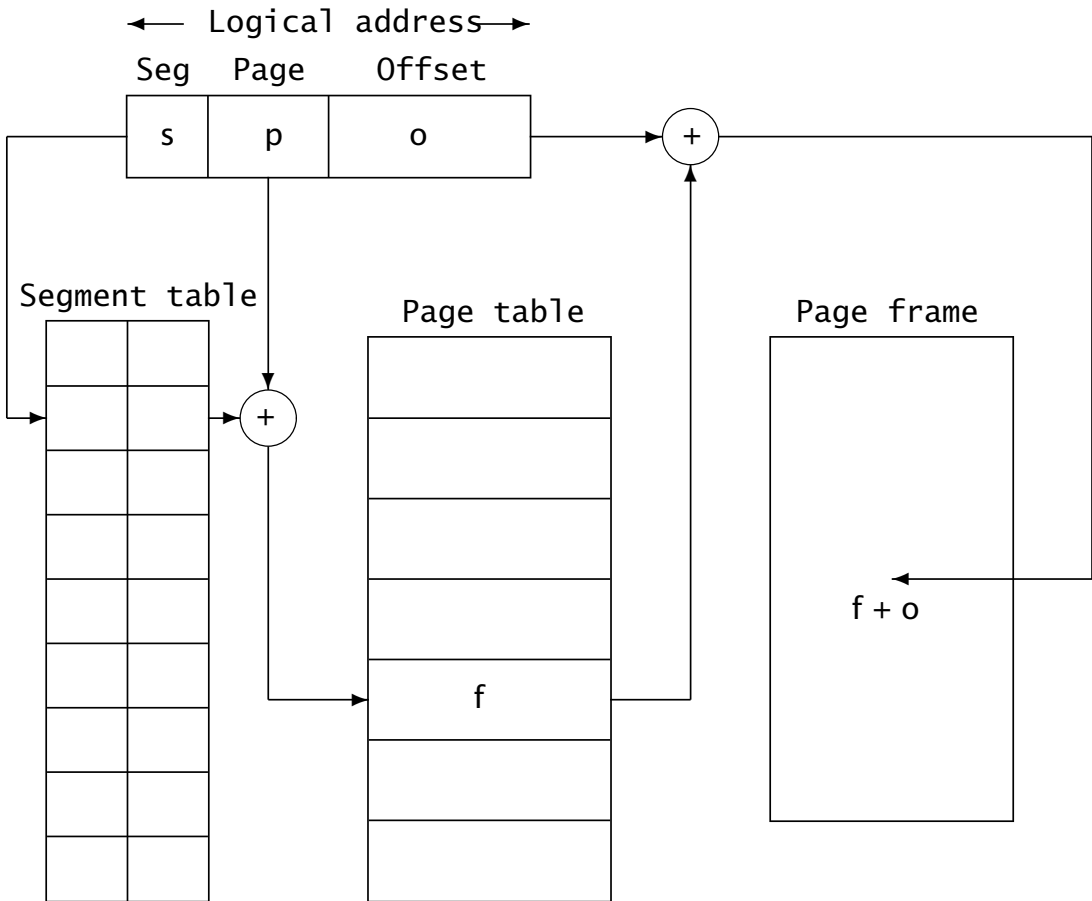
**Figure 8.29:** Effect of window size on working set



**Figure 8.30:** An address interpreted as segment, page, offset


```
Split the program address into s, p, o
Use s to index into the segment table
IF the presence bit in the segment descriptor is cleared
    (This means there are no pages for this segment in memory,
    nor is there a page table for this segment)
THEN create a new (empty) page table for this segment,
    enter its address in the segment table,
    and set the corresponding presence bit
ENDIF
Extract the address of the page table
Use p to index the page table
IF the presence bit in the corresponding entry
    in the page table is cleared (This means that
    the page is not in memory)
THEN fetch the page from backing store,
    enter its address in the pagetable,
    and set the corresponding presence bit
ENDIF
Extract the page frame address f
Add f to o to give the required location
```

**Figure 8.31:** Algorithm for paged segmentation

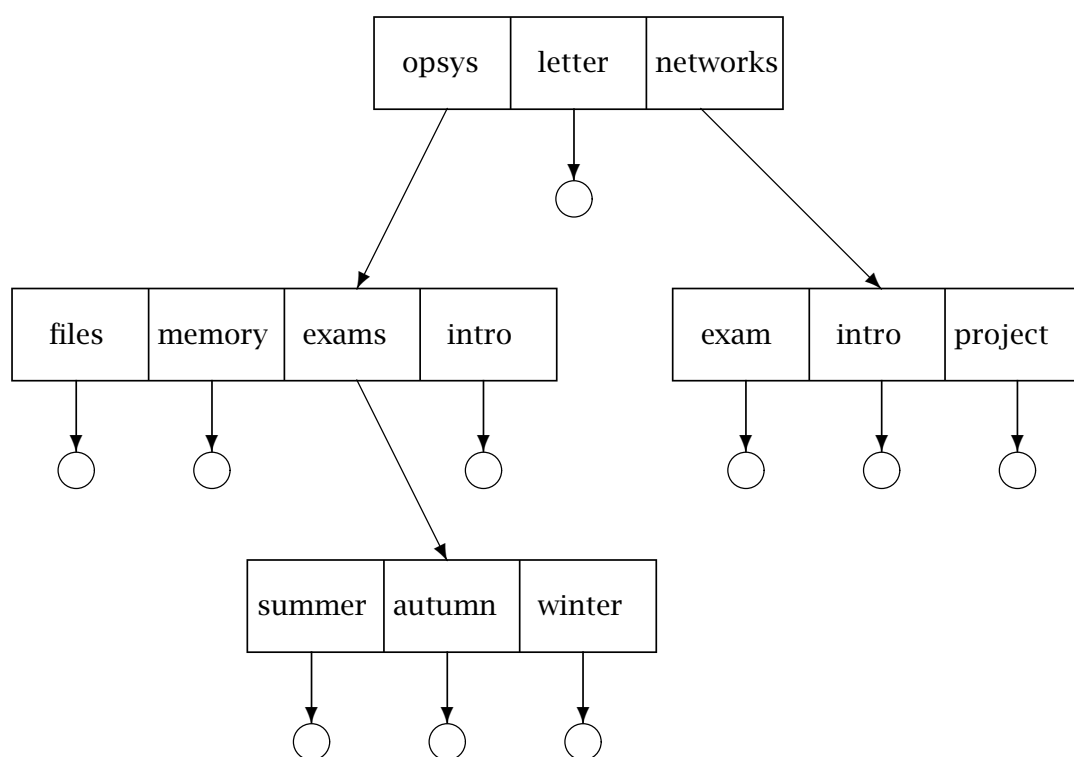


**Figure 8.32:** Address translation with paged segmentation





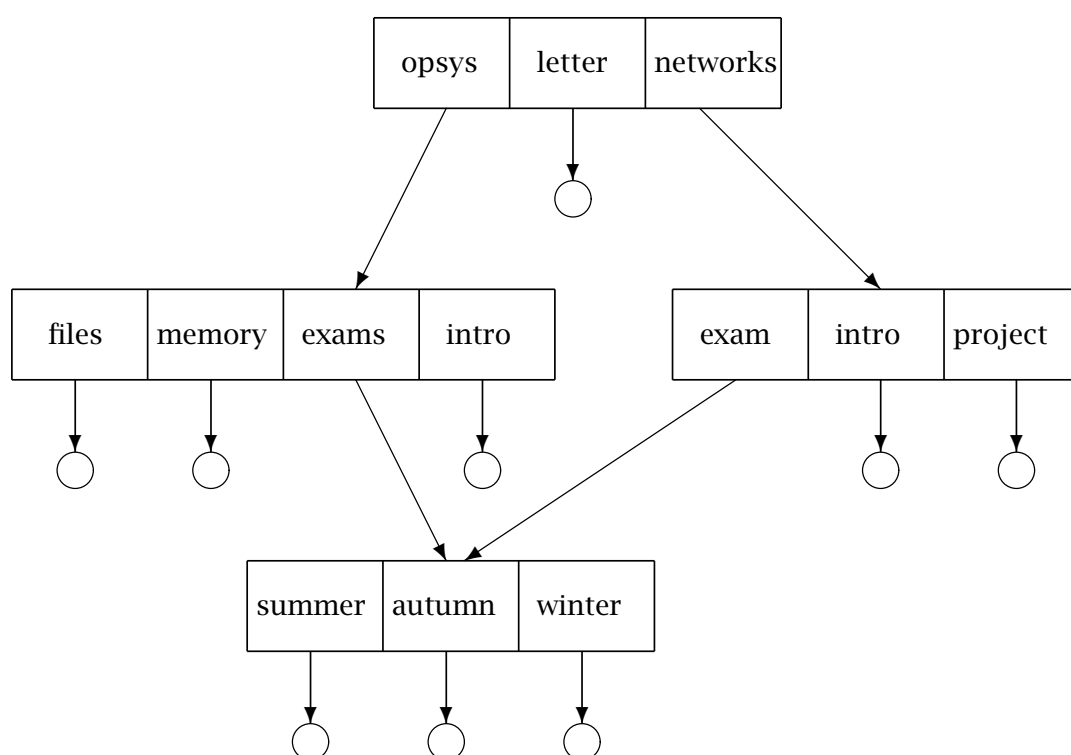
**Figure 9.1:** Overview of I/O processing



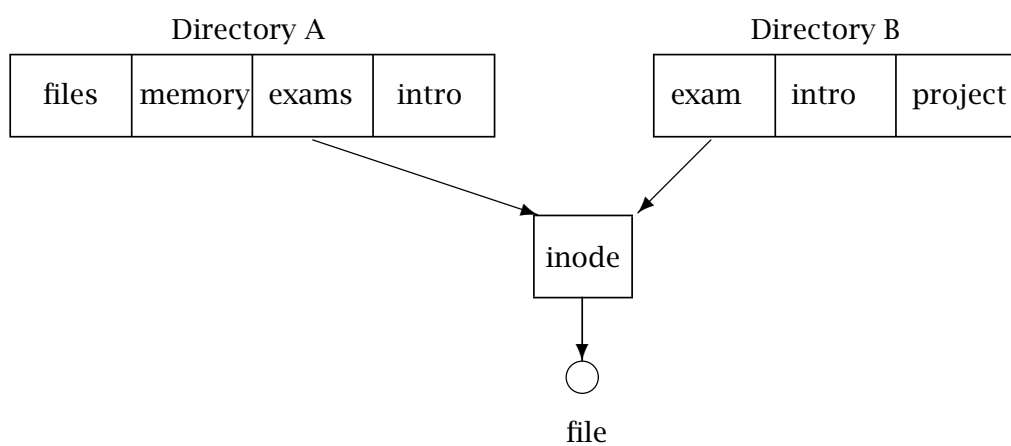
**Figure 9.2:** A tree-structured directory

```
struct dirent{
    long int      d_ino;    /* file number of this entry */
    unsigned short int d_namlen; /* length of string in d_name */
    unsigned short int d_reclen; /* length of this record */
    char          d_name[]; /* actual filename */
};
```

**Figure 9.3:** Structure of a directory entry



**Figure 9.4:** A directory with links



**Figure 9.5:** How a hard link is implemented

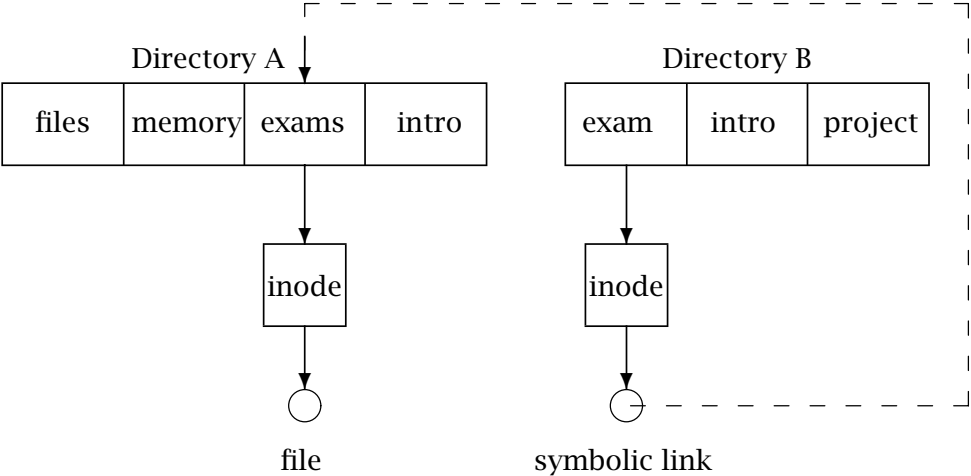


Figure 9.6: How a symbolic link is implemented

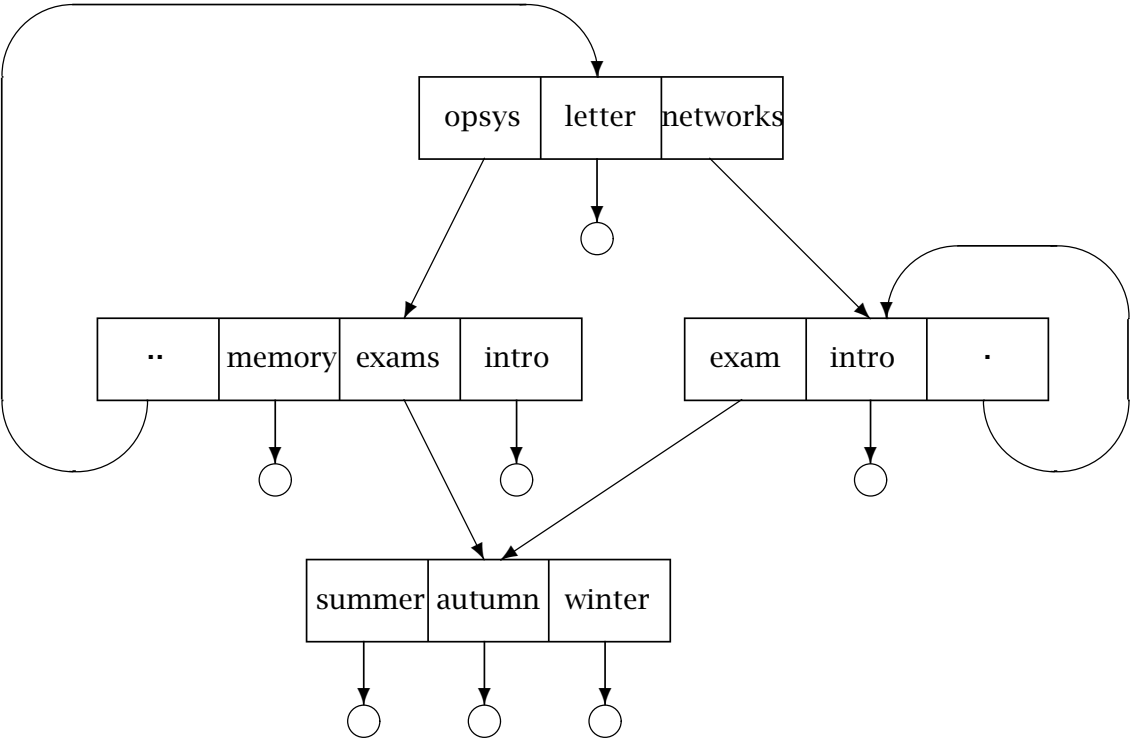


Figure 9.7: A directory structure with cycles

```

struct inode{
    struct list_head    i_hash;    /* link on hash list    */
    struct list_head    i_list;    /* link on inode list  */
    unsigned long       i_ino;     /* inode number        */
    unsigned short      i_count;   /* count of users      */
    kdev_t              i_dev;     /* device number       */
    umode_t             i_mode;    /* type and permissions */
    nlink_t             i_nlink;   /* hard links          */
    uid_t               i_uid;     /* owner               */
    gid_t               i_gid;     /* owner's group       */
    kdev_t              i_rdev;    /* if a device file    */
    off_t               i_size;    /* in bytes            */
    time_t              i_atime;   /* access              */
    time_t              i_mtime;   /* change of contents  */
    time_t              i_ctime;   /* inode modified      */
    unsigned long       i_blksize; /* size of block       */
    unsigned long       i_blocks;  /* number of blocks    */
    struct semaphore     i_sem;    /* mutex on inode      */
    struct inode_operations *i_op; /* operations vector    */
    struct super_block   *i_sb;    /* if a mount point    */
    struct wait_queue    *i_wait;  /* wait queue          */
    struct file_lock     *i_flock; /* locks on this file  */
    struct vm_area_struct *i_mmap; /* memory region       */
    struct page          *i_pages; /* pages in memory     */
    union{
        struct pipe_inode_info  pipe_i;
        struct ext2_inode_info  ext2_i;
        struct msdos_inode_info msdos_i;
        struct nfs_inode_info   nfs_i;
        struct socket           socket_i;
    }u;
};

```

**Figure 9.8:** Inode for an open file






Figure 9.9: Inodes on LRU and hash lists

```
#define S_IFSOCK 0140000 /* socket */
#define S_IFLNK 0120000 /* symbolic link */
#define S_IFREG 0100000 /* regular */
#define S_IFBLK 0060000 /* block special */
#define S_IFDIR 0040000 /* directory */
#define S_IFCHR 0020000 /* character special */
#define S_IFIFO 0010000 /* fifo */
#define S_ISUID 0004000 /* set user id on execution */
#define S_ISGID 0002000 /* set group id on execution */
#define S_IRUSR 0000400 /* read permission: owner */
#define S_IWUSR 0000200 /* write permission: owner */
#define S_IXUSR 0000100 /* execute/search permission: owner */
#define S_IRGRP 0000040 /* read permission: group */
#define S_IWGRP 0000020 /* write permission: group */
#define S_IXGRP 0000010 /* execute/search permission: group */
#define S_IROTH 0000004 /* read permission: other */
#define S_IWOTH 0000002 /* write permission: other */
#define S_IXOTH 0000001 /* execute/search permission: other */
```

Figure 9.10: Mode bit values

```
struct inode_operations{
    struct file_operations *default_file_ops;
    int (*create) ();
    struct dentry (*lookup) ();
    int (*link) ();
    int (*unlink) ();
    int (*symlink) ();
    int (*mkdir) ();
    int (*rmdir) ();
    int (*mknod) ();
    int (*rename) ();
    int (*readlink) ();
    struct dentry (*followlink) ();
    int (*readpage) ();
    int (*writepage) ();
    int (*bmap) ();
};
```

**Figure 9.11:** Some functions of the inode interface

```
struct dentry{
    struct inode    *d_inode;
    struct list_head d_hash;
    struct list_head d_lru;
    unsigned char    d_iname[];
};
```

**Figure 9.12:** An element in the cache of recent lookups

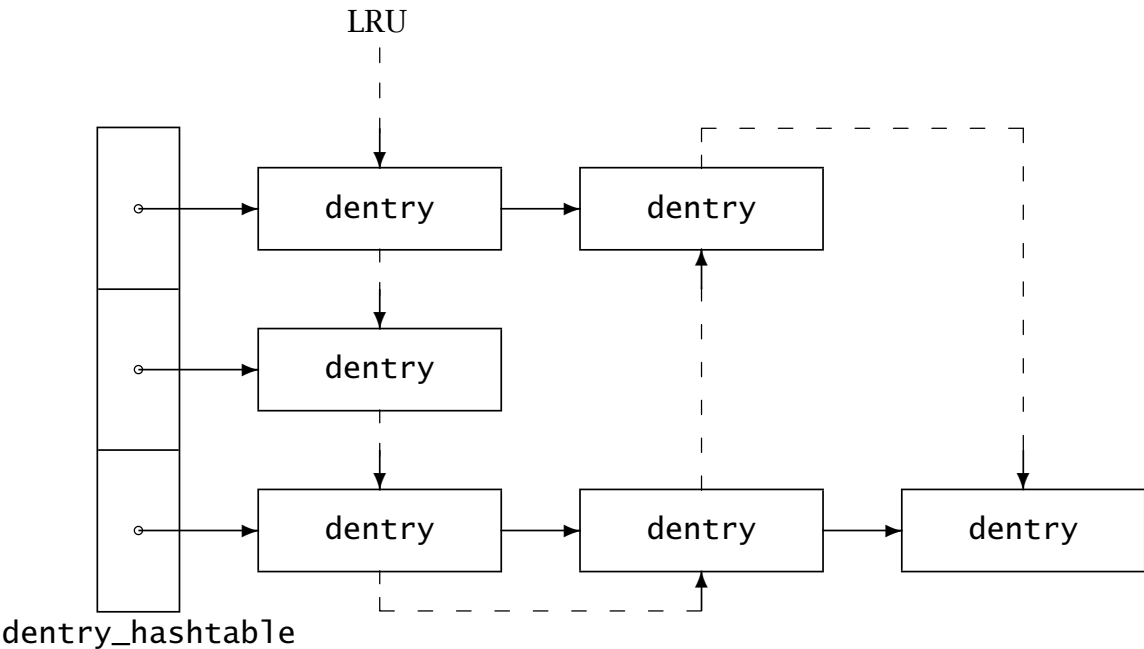


Figure 9.13: Directory name lookup cache

Owner	Group	Other
r w e	r w e	r w e

**Figure 9.14:** A protection bitmap

```
struct files_struct{
    atomic_t    count;
    fd_set      close_on_exec;
    fd_set      open_fds;
    struct file *fd_array[];
};
```

**Figure 9.15:** Definition of file descriptor table

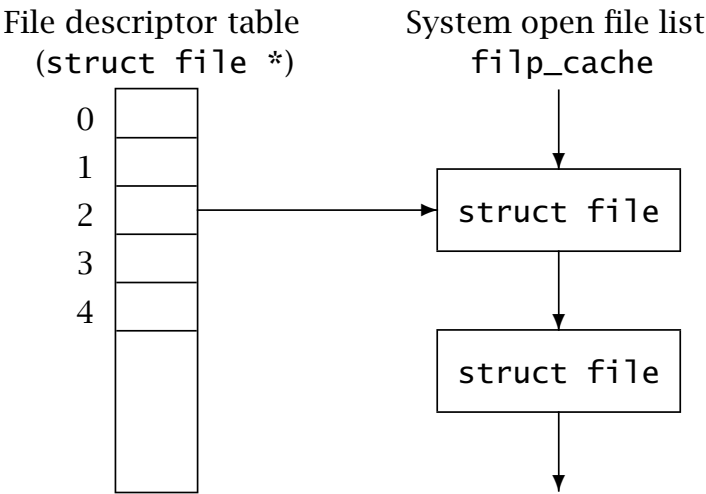


Figure 9.16: Descriptor and open file data structure



```
struct file{
    struct file      *f_next;    /* next struct file      */
    struct dentry    *f_dentry;  /* representing this stream */
    struct file_operations *f_op; /* operations on this stream */
    mode_t           f_mode;     /* open mode (read/write)   */
    loff_t           f_pos;      /* current position in file  */
    unsigned short   f_count;    /* number of threads sharing */
    int              f_owner;    /* where SIGIO should be sent */
};
```

**Figure 9.17:** Structure representing an open stream

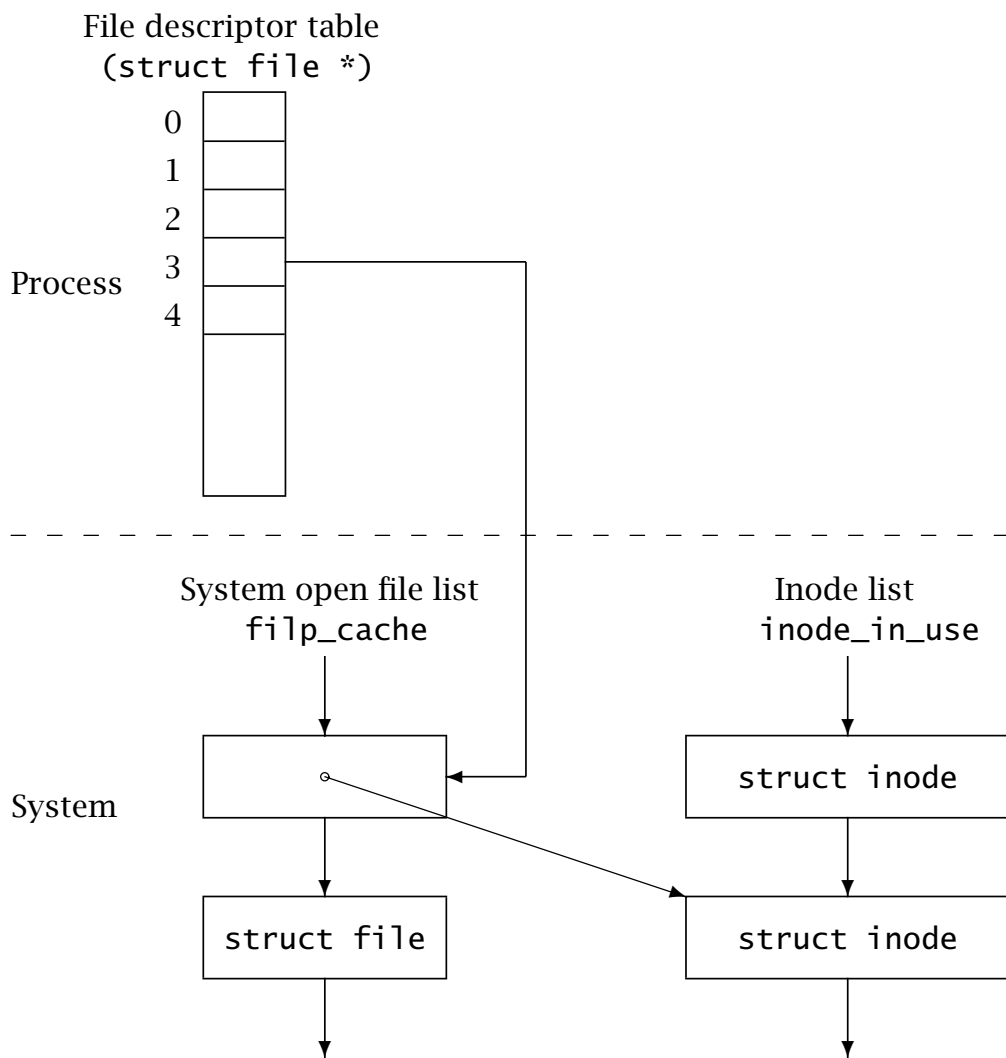


Figure 9.18: File opened by one user

```
struct file_operations{
    loff_t      (*llseek) ();
    ssize_t     (*read) ();
    ssize_t     (*write) ();
    int         (*readdir) ();
    unsigned int (*poll) ();
    int         (*ioctl) ();
    int         (*mmap) ();
    int         (*open) ();
    void        (*release)();
};
```

**Figure 9.19:** Operations on open files

```
Set up file descriptor, link to struct file
Search the directory name cache
IF not found THEN
    lookup() (from directory i_op)
    IF not found THEN
        errno = EFAULT
        Return (-1)
    ENDIF
    Create appropriate entry in name cache
    Create inode in memory
    Link dentry to inode
ENDIF
Set up f_op field in struct file
open() (from f_op)
IF no permission THEN
    errno = EACCES
    Return (-1)
ENDIF
Link struct file to dentry
Increment i_count
Return (file descriptor)
```

**Figure 9.20:** Algorithm for open

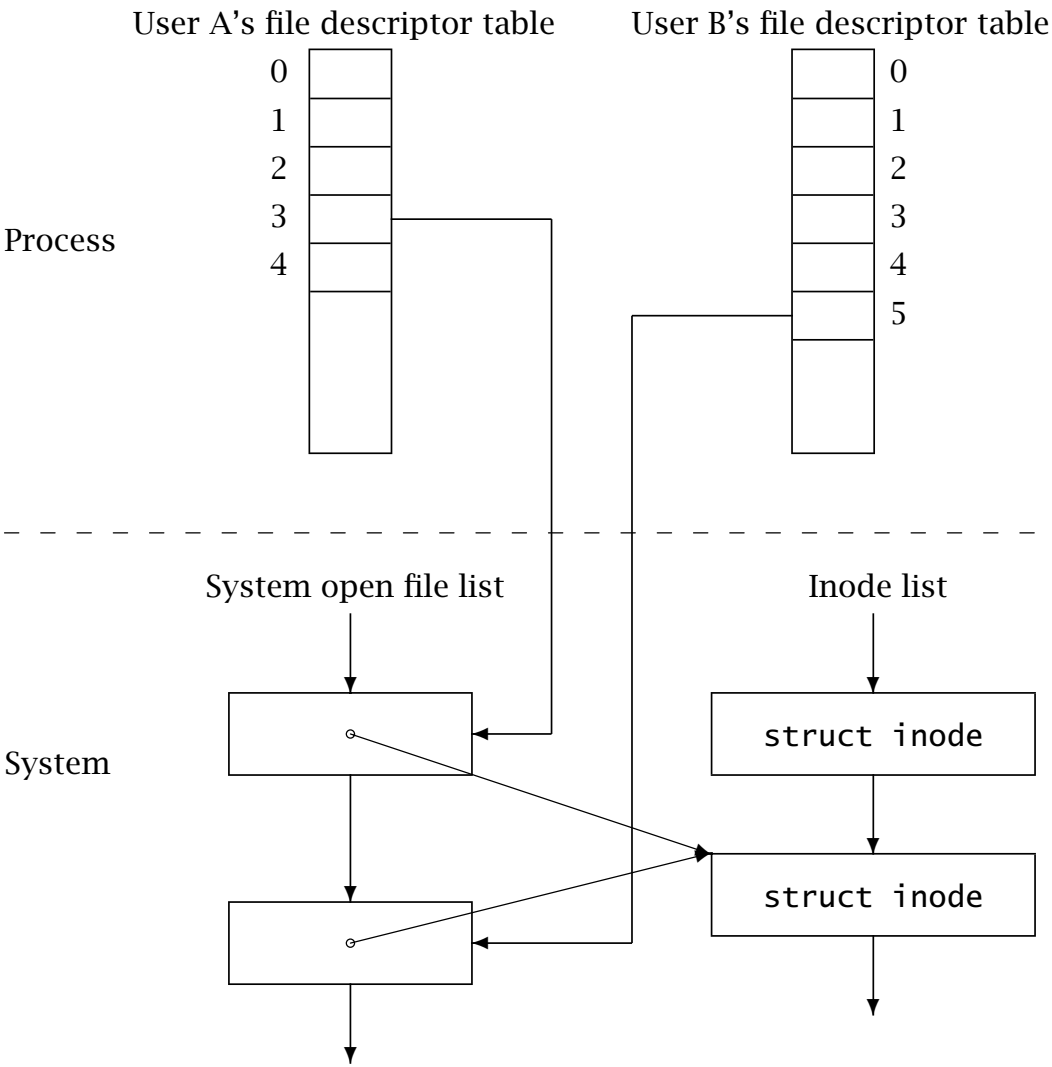


Figure 9.21: File opened by two users

```
Map fileid onto stream
IF not open, THEN
    errno = errornumber
    Return (-1)
ENDIF
Check mode and length against stream characteristics
IF not compatible THEN
    errno = errornumber
    Return (-1)
ENDIF
Call stream specific function with appropriate parameters
IF not successful THEN
    errno = errornumber
    Return (-1)
ENDIF
Return (0)
```

**Figure 9.22:** Generic I/O algorithm


```
Use stream number to index into the file descriptor table
IF not assigned, THEN
    errno = EBADF
    Return (-1)
ENDIF
Follow the pointer from descriptor to struct file
Examine f_mode
IF request not compatible with open mode THEN
    errno = EBADF
    Return (-1)
ENDIF
IF request is for simple I/O THEN
    Call appropriate function in f_op
ELSE
    Follow f_dentry pointer to struct inode
    Call appropriate function in i_op
ENDIF
IF not successful THEN
    errno = EIO
    Return (-1)
ELSE
    Return (0)
ENDIF
```

**Figure 9.23:** Generic algorithm for Unix I/O procedure

```
struct pollfd{
    int fd;           /* the descriptor we are interested in */
    short int events; /* flag specifying events of interest */
    short int revents; /* flag specifying events which have occurred */
};
```

**Figure 9.24:** Data structure for polling a descriptor





**Figure 9.25:** Several processes waiting on a device

```
struct file_lock{
    struct file_lock    *fl_next;
    struct task_struct  *fl_owner;
    struct wait_queue   *fl_wait;
    struct file          *fl_file;
    unsigned char        fl_type;
    off_t                fl_start;
    off_t                fl_end;
};
```

**Figure 9.26:** Structure representing a lock on a file




Figure 9.27: A lock request which blocks




Figure 9.28: How locks are recorded

```
Check parameters against device characteristics
Return if error
IF buffer_empty THEN
    Call device driver (for bufferfull)
    Process any errors
ENDIF
Transfer data from buffer
Return
```

**Figure 9.29:** A buffered input procedure

```
struct buffer_head{
    unsigned long    b_blocknr;    /* block number    */
    kdev_t           b_dev;        /* device          */
    struct buffer_head *b_next;    /* hash list       */
    unsigned long    b_state;      /* state bitmap    */
    struct buffer_head *b_next_free; /* free list       */
    char             *b_data;      /* pointer to data block */
    struct wait_queue *b_wait;     /* processes waiting */
};
```

**Figure 9.30:** A buffer head

```
#define BH_Uptodate 0 /* contains valid data */
#define BH_Dirty    1 /* is dirty */
#define BH_Lock     2 /* is locked */
#define BH_Req      3 /* has been invalidated */
```

**Figure 9.31:** Bits representing the state of a buffer

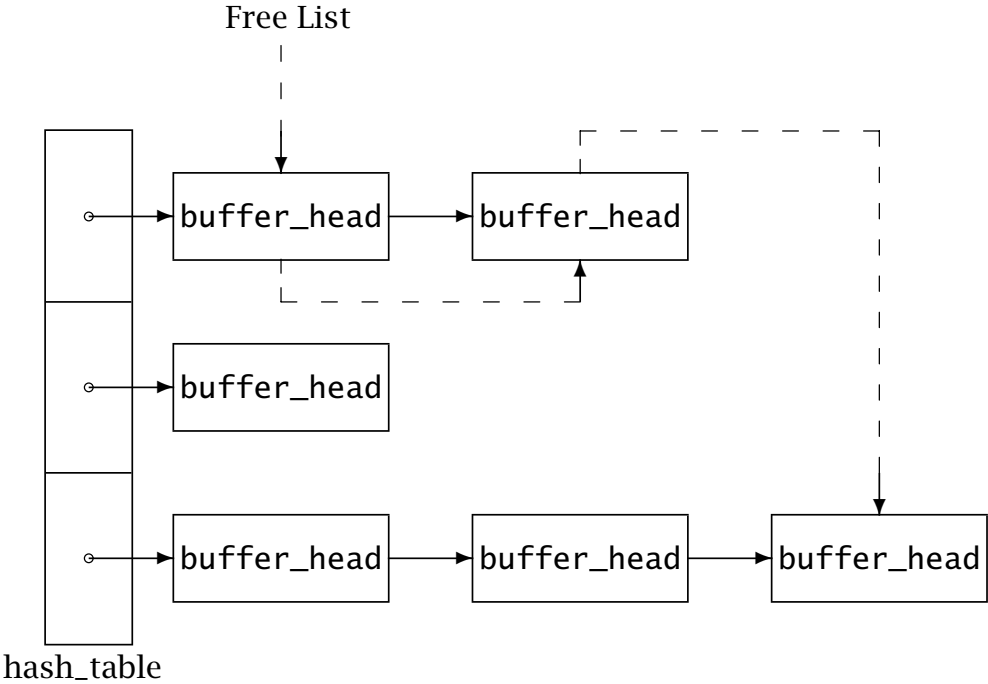
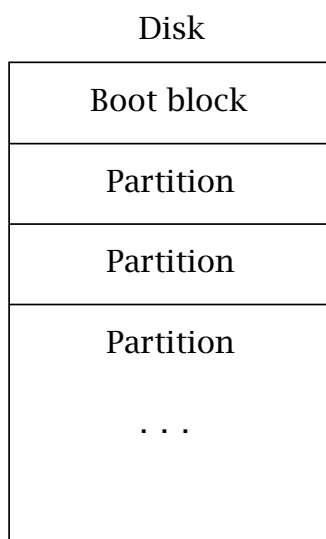


Figure 9.32: Buffer cache



```
struct file_system_type{
    const char          *name;
    struct super_block  *(*read_super)();
    struct file_system_type *next;
};
```

**Figure 10.1:** An entry in the linked list of file systems



**Figure 10.2:** Layout of a disk

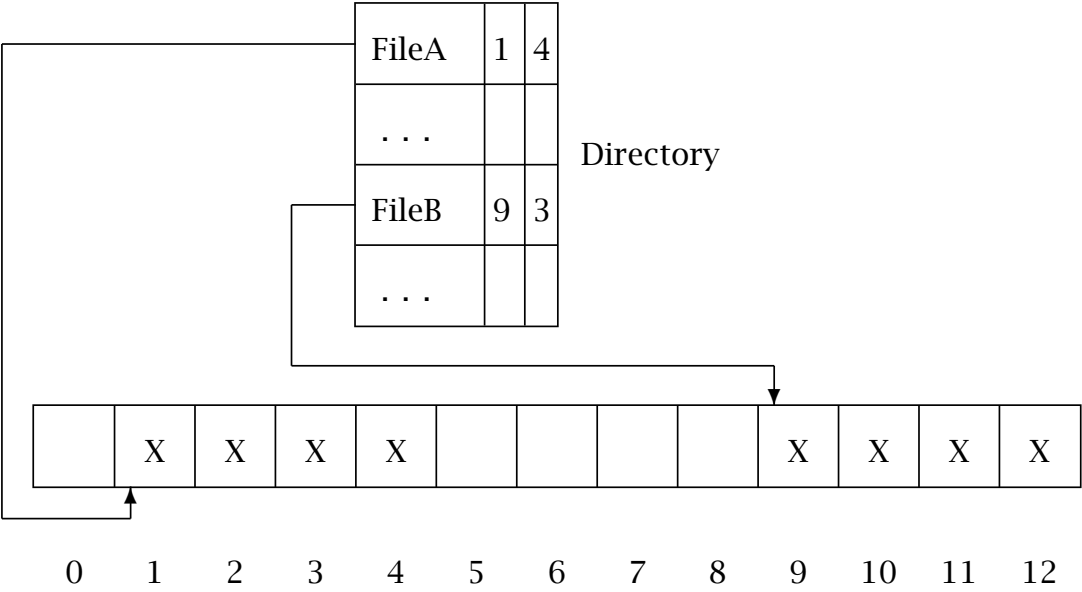


Figure 10.3: Contiguous allocation

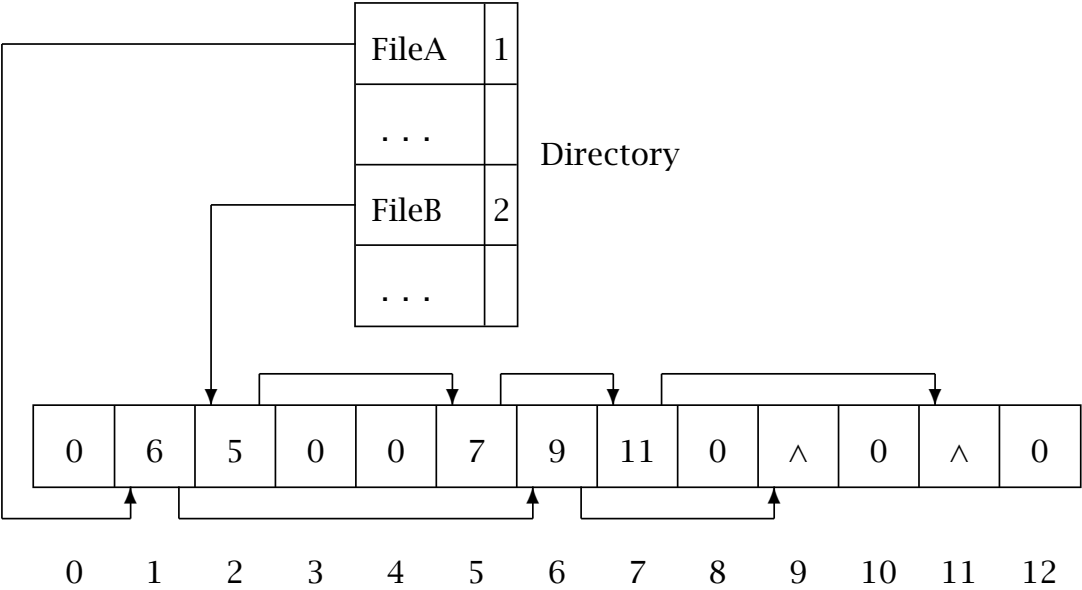


Figure 10.4: Linked allocation

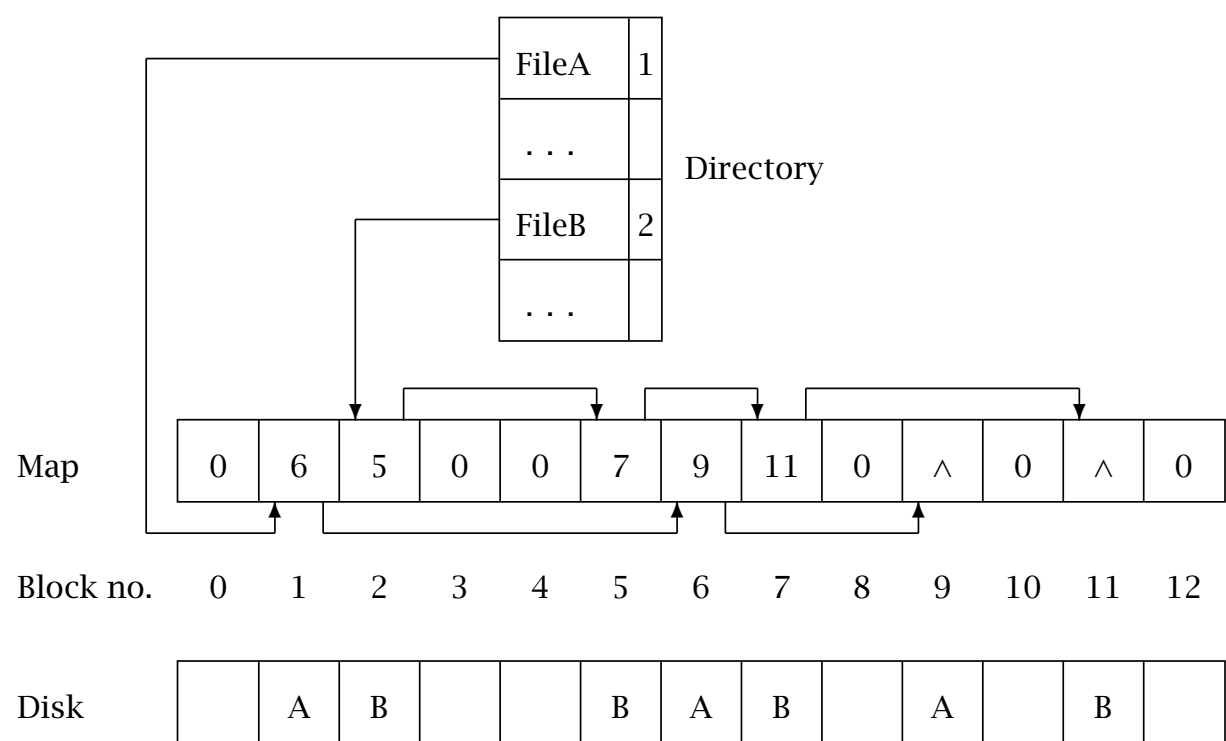


Figure 10.5: Example of a file map

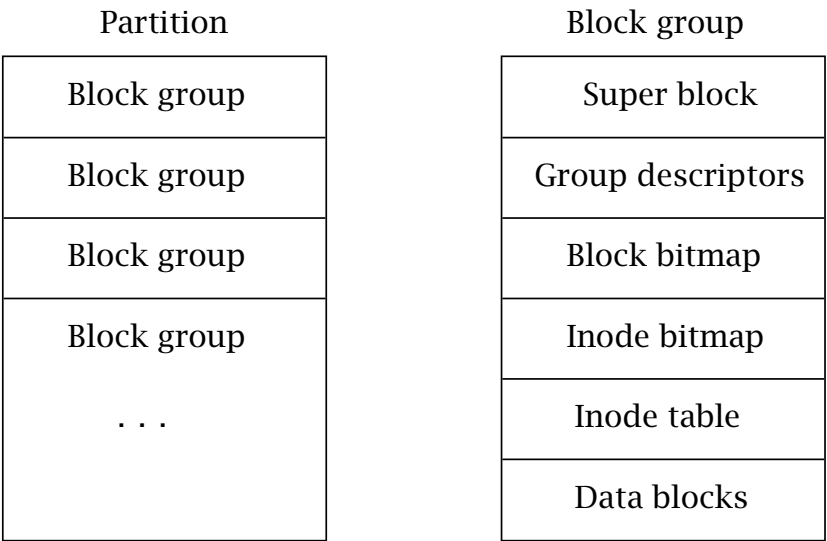



Figure 10.6: Layout of a partition with Ext2



**Figure 10.7:** A block group on disk

```
struct ext2_super_block{
    u32 s_inodes_count;      /* Inodes count          */
    u32 s_blocks_count;     /* Blocks count       */
    u32 s_free_blocks_count; /* Free blocks count  */
    u32 s_free_inodes_count; /* Free inodes count  */
    u32 s_first_data_block; /* First data block   */
    u32 s_log_block_size;   /* Block size         */
    s32 s_log_frag_size;    /* Fragment size      */
    u32 s_blocks_per_group; /* # Blocks per group */
    u32 s_inodes_per_group; /* # Inodes per group */
    u16 s_inode_size;       /* size of inode structure */
    u16 s_block_group_nr;   /* number of this block group */
};
```

**Figure 10.8:** An Ext2 super block



```
struct ext2_group_desc{
    u32 bg_block_bitmap;      /* Blocks bitmap block */
    u32 bg_inode_bitmap;      /* Inodes bitmap block */
    u32 bg_inode_table;       /* Inodes table block */
    u16 bg_free_blocks_count; /* Free blocks count */
    u16 bg_free_inodes_count; /* Free inodes count */
    u16 bg_used_dirs_count;   /* Directories count */
};
```

**Figure 10.9:** Structure of a block group descriptor

```
struct ext2_inode{
    u16 i_mode;           /* File mode           */
    u16 i_uid;            /* Owner uid           */
    u32 i_size;           /* Size in bytes       */
    u32 i_atime;          /* Access time         */
    u32 i_ctime;          /* Creation time       */
    u32 i_mtime;          /* Modification time   */
    u16 i_gid;            /* Group id            */
    u16 i_links_count;    /* Links count         */
    u32 i_blocks;         /* Blocks count        */
    u32 i_block[15];      /* Pointers to blocks  */
    u32 i_faddr;          /* Fragment address    */
    u8  l_i_frag;          /* Fragment number     */
    u8  l_i_fsize;        /* Fragment size       */
};
```

**Figure 10.10:** Structure of an Ext2 disk inode

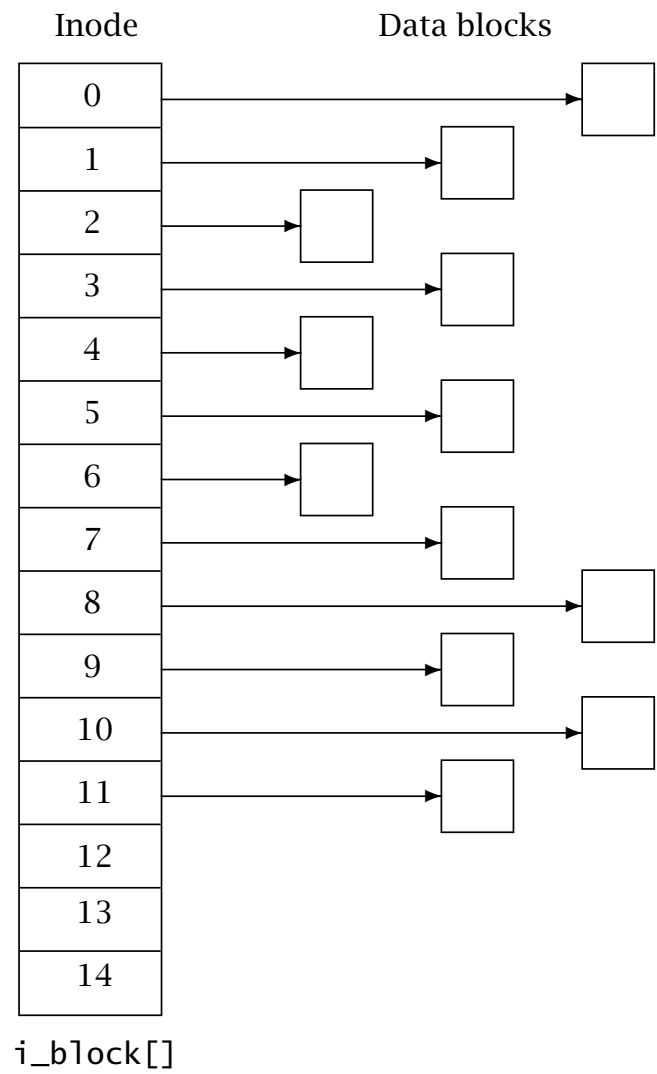



Figure 10.11: Allocation information in an inode



**Figure 10.12:** Single indirect block




Figure 10.13: Double indirect block




Figure 10.14: Fragments of two files in a block

```
struct ext2_dir_entry{
    u32  inode;    /* inode number of the entry */
    u16  rec_len;  /* total length of entry      */
    u16  name_len; /* length of name          */
    char name[]    /* file name                */
};
```

**Figure 10.15:** An Ext2 directory entry

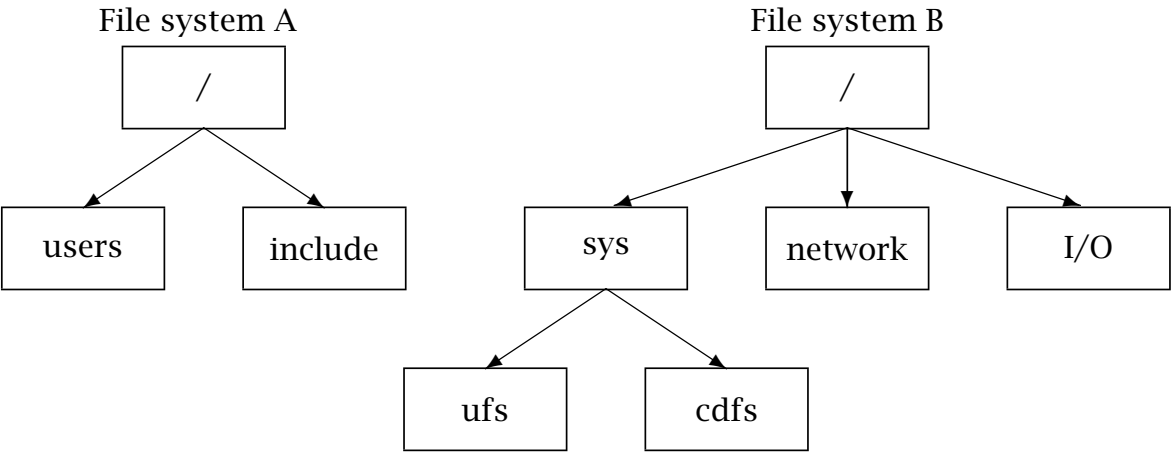



Figure 10.16: File systems as they exist on disk





**Figure 10.17:** Directory structure after attaching file system B

```
struct super_block{
    struct list_head      s_list;
    kdev_t                s_dev;
    struct file_system_type *s_type;
    struct super_operations *s_op;
    struct dentry          *s_root;
    union{
        struct ext2_sb_info  ext2_sb;
        struct msdos_sb_info msdos_sb;
        struct nfs_sb_info   nfs_sb;
    }u;
};
```

**Figure 10.18:** Data structure representing a mounted file system

```
struct super_operations{
    void (*read_inode)();
    void (*write_inode)();
    void (*put_inode)();
    void (*put_super)();
    void (*write_super)();
    int  (*statfs)();
};
```

**Figure 10.19:** Operations on a file system

```
struct ext2_sb_info{
    struct ext2_super_block *s_es;
    struct buffer_head      **s_group_desc;
    struct buffer_head      *s_inode_bitmap[];
    struct buffer_head      *s_block_bitmap[];
};
```

**Figure 10.20:** Superblock for an Ext2 file system

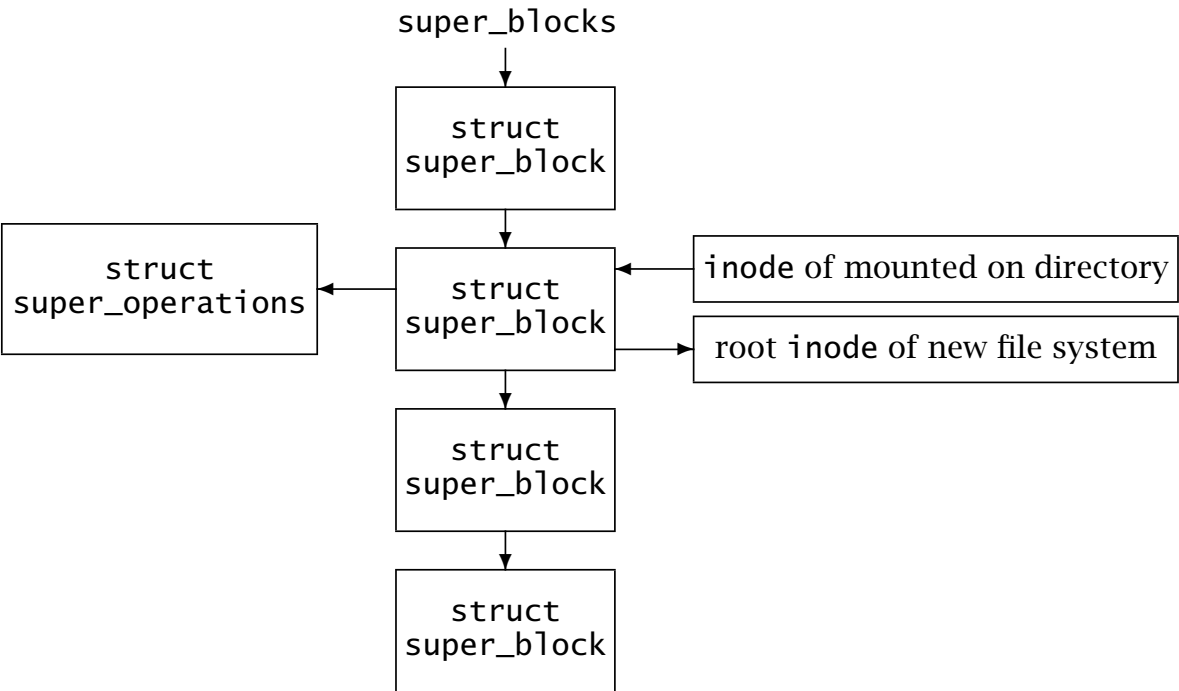


Figure 10.21: Data structures involved in mounting a file system

```
struct ext2_inode_info{
    u32 i_data[15];
    u32 i_faddr;
    u8  i_frag_no;
    u8  i_frag_size;
    u32 i_file_acl;
    u32 i_dir_acl;
    u32 i_block_group;
};
```

**Figure 10.22:** Ext2 file system inode data in memory

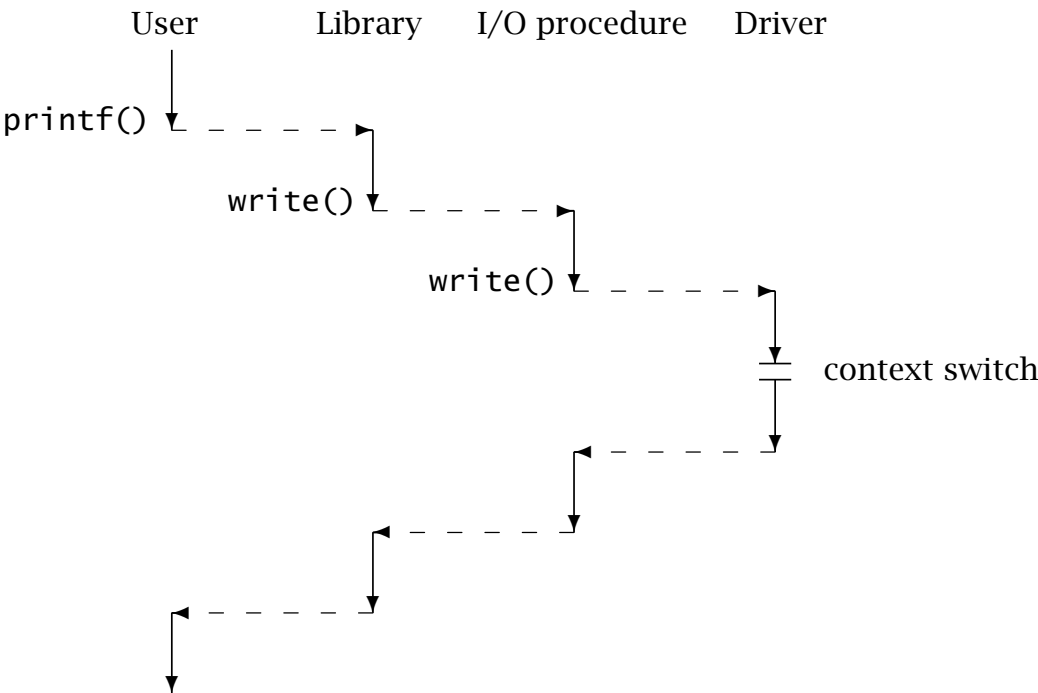


Figure 11.1: Calling a Unix style driver

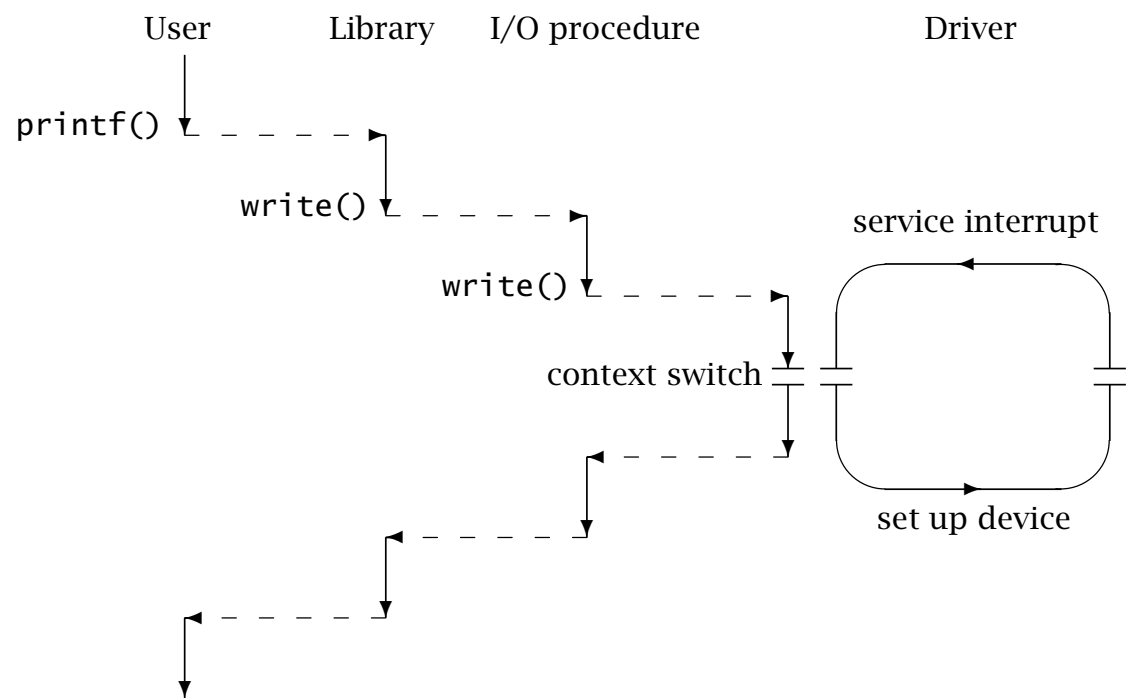



Figure 11.2: Device driver as a process



```
struct device_struct{
    const char          *name;
    struct file_operations *fops;
};

static struct device_struct chrdevs[];
static struct device_struct blkdevs[];
```

**Figure 11.3:** The character and block device switches



**Figure 11.4:** An entry in the character device switch

```
IF device busy THEN
    Sleep
ENDIF
Set up the I/O operation
Sleep
IF error THEN
    return value = errornumber
ELSE
    Transfer data to destination
    return value = success
ENDIF
Return (return value)
```

**Figure 11.5:** Algorithm for reading from a character device

```
struct request{
    kdev_t      rq_dev;      /* device identifier */
    int          cmd;         /* READ or WRITE      */
    unsigned long sector;    /* start sector        */
    unsigned long nr_sectors; /* number of sectors   */
    char         *buffer;    /* buffer address      */
    struct request *next;    /* next IORB           */
};
```

**Figure 11.6:** Request block for block devices

```
Put struct request on device queue
IF device idle, THEN
    Select struct request from queue
    IF (cmd == WRITE) THEN
        Transfer data from user space to device memory
    ENDIF
    Set up the device for specified operation
ENDIF
Return to caller.
```


**Figure 11.7:** Algorithm for strategy routine

```
Check for error
IF (cmd == READ) THEN
    Transfer data from device memory to user space
ENDIF
Move user process from wait queue to run queue
IF request queue not empty THEN
    Select struct request from queue
    IF (cmd == WRITE) THEN
        Transfer data from user space to device memory
    ENDIF
    Set up the device for specified operation
ENDIF
```

**Figure 11.8:** Algorithm for interrupt service routine

```
IF device busy THEN
    Sleep
ENDIF
Transfer data from user buffer to device memory
Set up the device for writing
Sleep
IF error THEN
    return value = errornumber
ELSE
    return value = success
ENDIF
Return (return value)
```

**Figure 11.9:** Algorithm for writing to a character device



**Figure 11.10:** A basic view of a stream









Figure 11.11: Data structures constituting a minimal stream



**Figure 11.12:** A many-to-one multiplexor



**Figure 11.13:** A one-to-many multiplexor



**Figure 11.14:** A many-to-many multiplexor




Figure 11.15: Networking multiplexor

```
struct termios{
    tcflag_t c_iflag;
    tcflag_t c_oflag;
    tcflag_t c_cflag;
    tcflag_t c_lflag;
    cc_t      c_cc[];
};
```

**Figure 11.16:** Data structure controlling terminal characteristics

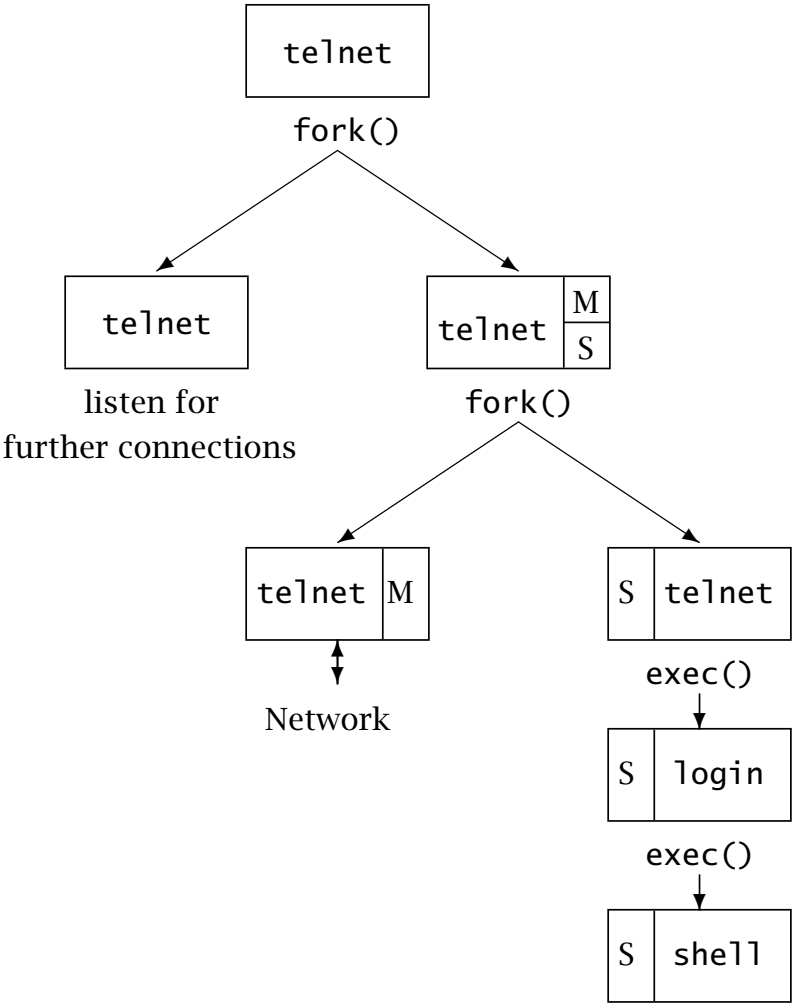
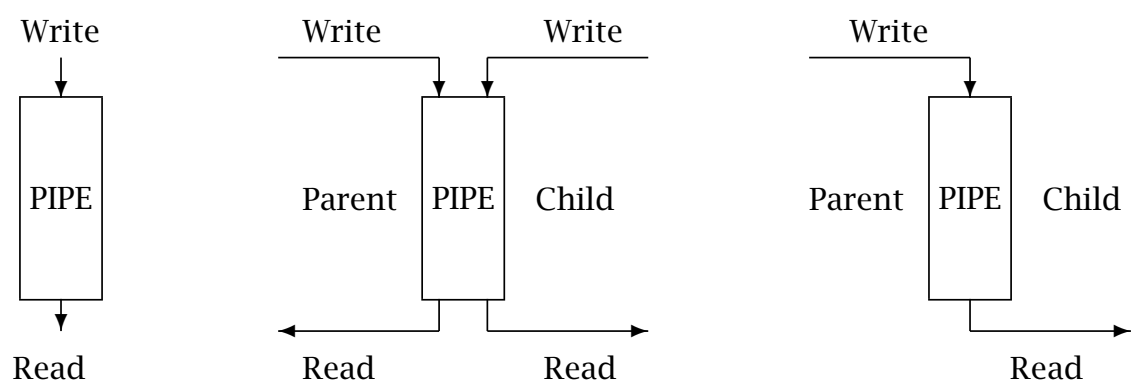


Figure 11.17: Logging on using a pseudo terminal








**Figure 12.1:** Setting up a pipe for interprocess communication

```
struct pipe_inode_info{
    struct wait_queue *wait    /* for blocked threads    */
    char *base    /* buffer    */
    unsigned int start    /* next byte to read    */
    unsigned int readers    /* number of readers    */
    unsigned int writers    /* number of writers    */
};
```

**Figure 12.2:** Data structure representing a pipe



**Figure 12.3:** Representation of a pipe

```
#define AF_UNSPEC    0 /* Unspecified          */
#define AF_UNIX      1 /* Unix domain sockets */
#define AF_INET      2 /* Internet IP Protocol */
#define AF_IPX       4 /* Novell IPX           */
#define AF_APPLETALK 5 /* Appletalk            */
```

**Figure 12.4:** Supported address families

```
struct net_proto_family{  
    int family;  
    int (*create)();  
};
```

**Figure 12.5:** Data structure representing a protocol family




Figure 12.6: Domains registered in the kernel

```
struct socket{
    socket_state      state; /* current state of connection */
    struct proto_ops  *ops;  /* domain specific functions */
    struct sock       *data; /* domain specific data */
    struct wait_queue **wait; /* waiting processes */
    short             type;  /* SOCK_STREAM, etc. */
};
```

**Figure 12.7:** Data structure representing a socket

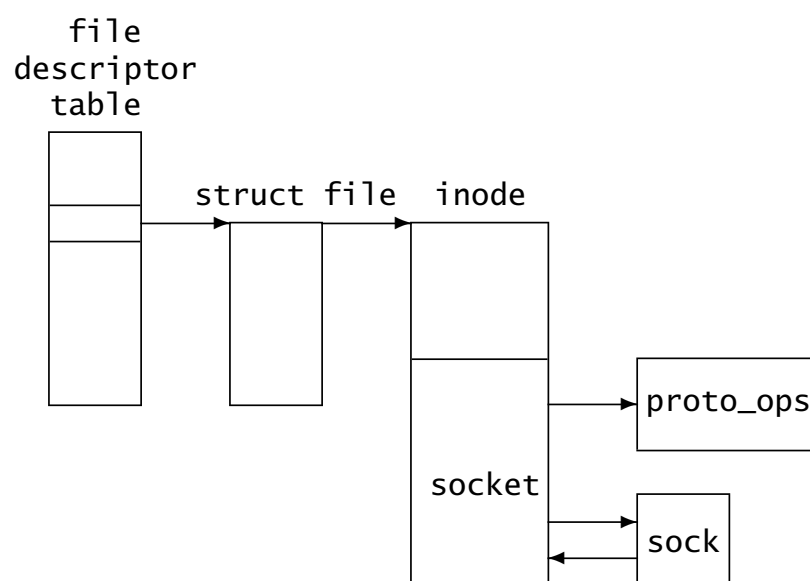
```
typedef enum{
    SS_FREE,          /* not allocated          */
    SS_UNCONNECTED,   /* unconnected to any socket */
    SS_CONNECTING,    /* in process of connecting */
    SS_CONNECTED,     /* connected to socket      */
    SS_DISCONNECTING  /* in process of disconnecting */
}socket_state;
```

**Figure 12.8:** State values for a socket



```
struct proto_ops{
    int family;
    int (*bind)();
    int (*connect)();
    int (*socketpair)();
    int (*accept)();
    int (*listen)();
    int (*shutdown)();
    int (*setsockopt)();
    int (*getsockopt)();
    int (*sendmsg)();
    int (*recvmsg)();
};
```

**Figure 12.9:** Data structure representing a protocol



**Figure 12.10:** Data structures after socket is allocated

```
struct sockaddr{
    sa_family_t sa_family; /* address family, AF_xxx */
    char        sa_data[]; /* protocol address      */
};
```


**Figure 12.11:** Generic format of a socket address

```
struct msghdr{
    void          *msg_name;   /* name of destination socket */
    int           msg_namelen; /* length of name             */
    struct iovec  *msg_iov;     /* array of data buffers      */
    int           msg_iovlen;   /* number of buffers          */
};
```

**Figure 12.12:** Message header

```
struct sock{
    struct sock      *next;
    int              rcvbuf;
    struct sk_buff_head receive_queue;
    int              sndbuf;
    struct sk_buff_head write_queue;
    struct proto      *prot;
    struct unix_opt   af_unix;
    struct socket     *socket;
};
```

**Figure 12.13:** Domain-specific control block



**Figure 12.14:** Hash chain of Unix domain control blocks

```
struct proto{
    void      (*close)();
    int       (*connect)();
    struct sock* (*accept) ();
    int       (*poll)();
    void      (*shutdown)();
    int       (*setsockopt)();
    int       (*getsockopt)();
    int       (*sendmsg)();
    int       (*recvmsg)();
    int       (*bind)();
};
```

**Figure 12.15:** A protocol specific structure

```
struct unix_opt{  
    struct unix_address *addr;  
    struct sock         *other;  
};
```

**Figure 12.16:** Links for a Unix specific socket



```
struct sockaddr_un{
    unsigned short sun_family;    /* address family, AF_UNIX */
    char           sun_path[104]; /* pathname                */
};
```

**Figure 12.17:** Structure containing address of a Unix domain socket




Figure 12.18: A pair of connected sockets

```
struct sk_buff{
    struct sk_buf *next;    /* Next buffer in list */
    struct sock *sk;        /* Socket we are owned by */
    unsigned long len;      /* Length of actual data */
    unsigned int truesize;   /* Buffer size */
    unsigned char *head;    /* Head of buffer */
    unsigned char *data      /* Data head pointer */
    unsigned char *tail     /* Tail pointer */
    unsigned char *end       /* End pointer */
};
```

**Figure 12.19:** Buffer used for communication over sockets






Figure 12.20: Data part of an `sk_buff`



**Figure 13.1:** Networked computers



**Figure 13.2:** Workstations connected on a LAN




Figure 13.3: Diskless workstations on a LAN




Figure 13.4: Processor pool model



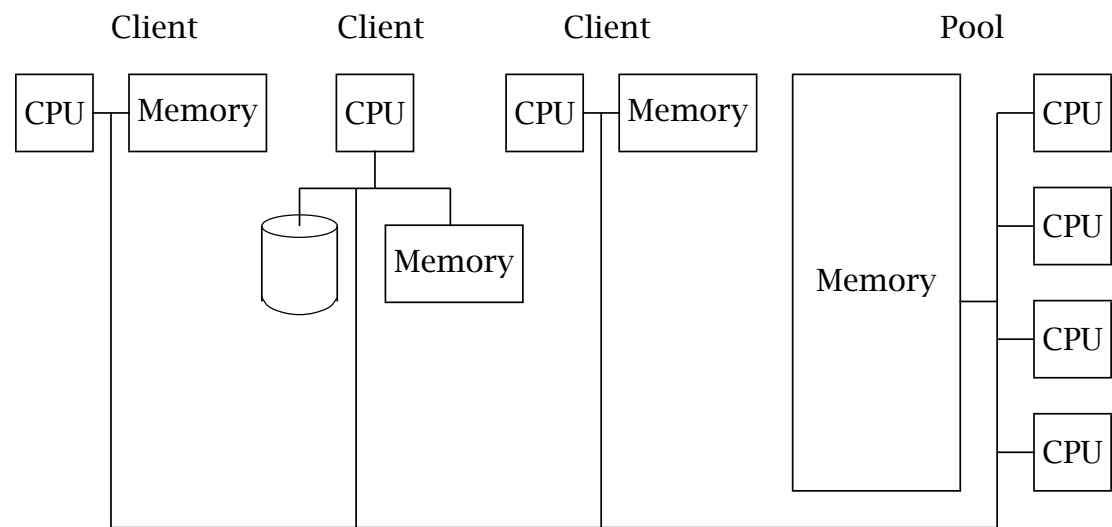

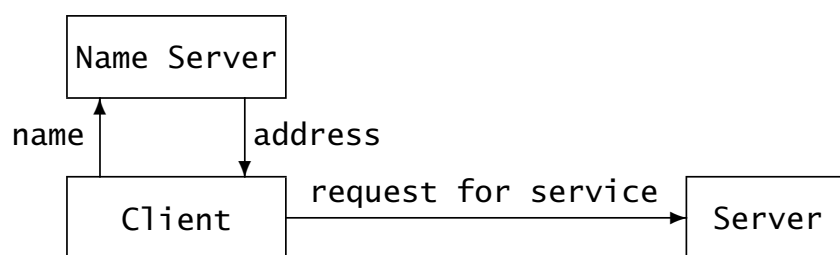


Figure 13.5: Combined model



**Figure 13.6:** Part of the Internet name space



**Figure 13.7:** Using a name server

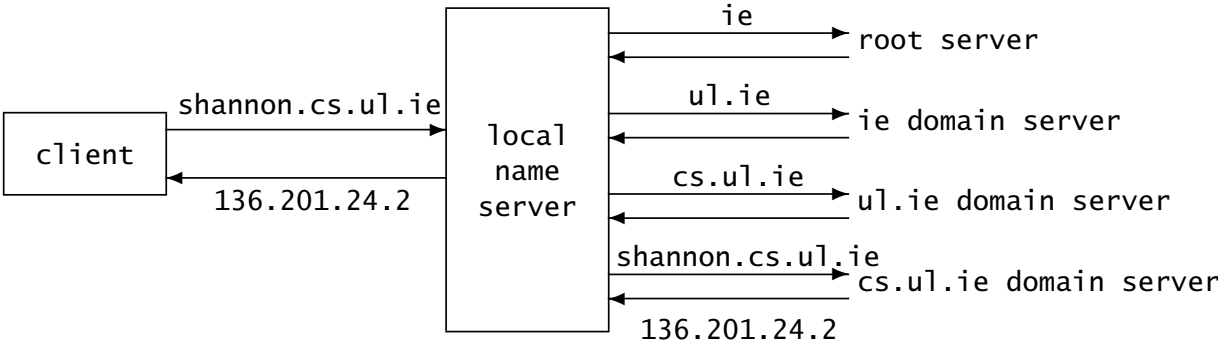



Figure 13.8: Translating an Internet name



**Figure 13.9:** Overview of CORBA

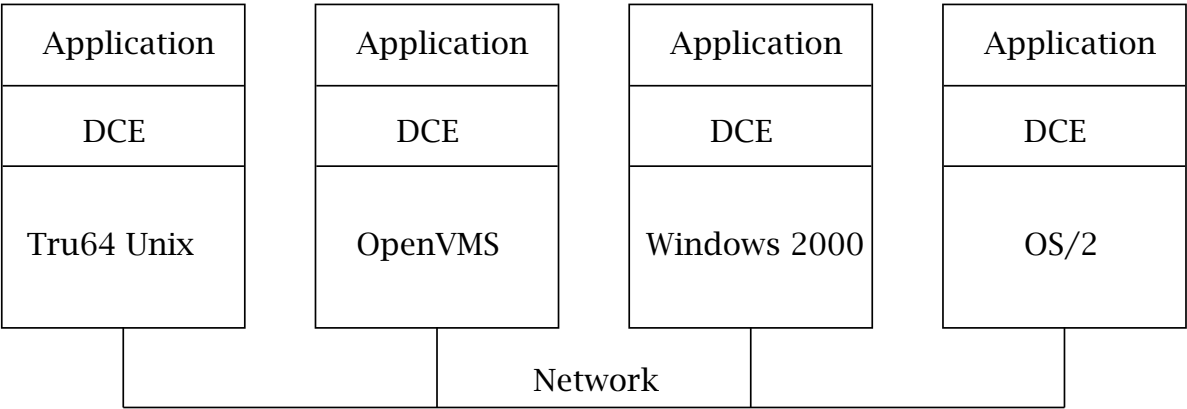


Figure 13.10: The distributed computing environment

```
struct device{
    char          *name;
    struct device *next;
    int           (*init)();
    unsigned short type;
    unsigned char dev_addr[];
    int           (*hard_start_xmit)();
};
```

**Figure 14.1:** Data structure describing a network interface

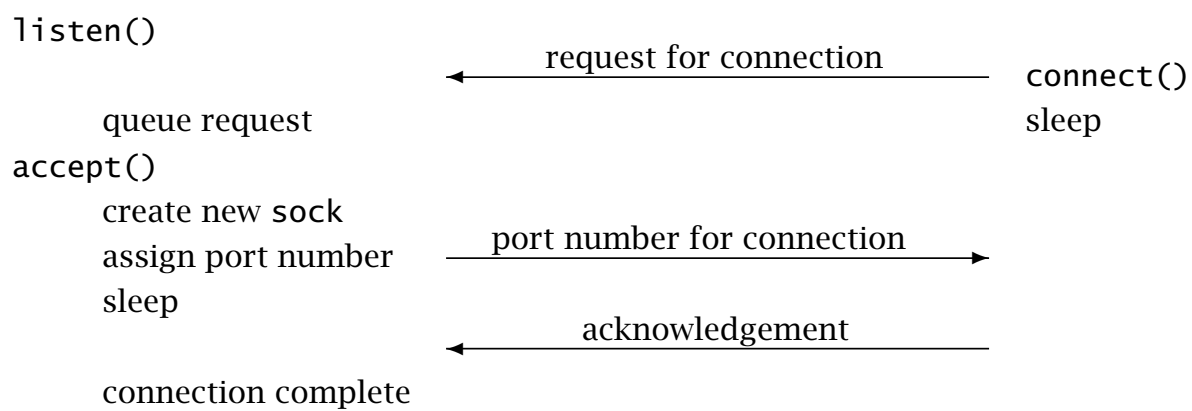
```
struct sock{
    struct options      *opt;
    struct sock         *next;
    struct sk_buff      *partial;
    struct sk_buff_head receive_queue, write_queue;
    struct proto        *prot;
    u32                 daddr;
    u32                 rcv_saddr;
    u16                 dport;
    unsigned short      num;
    u32                 saddr;
    struct socket       *socket;
    void                (*data_ready)();
};
```

**Figure 14.2:** An Internet protocol control block




```
sockaddr_in{
    short int      sin_family; /* Address family, AF_INET */
    unsigned short int sin_port; /* Port number */
    struct in_addr  sin_addr; /* Internet address */
};
```

**Figure 14.3:** Internet socket address structure



**Figure 14.4:** Connecting two sockets



**Figure 14.5:** A physical frame

```
struct ethhdr{  
    unsigned char  h_dest[];  
    unsigned char  h_source[];  
    unsigned short h_proto;  
};
```

**Figure 14.6:** An ethernet header

```
struct packet_type{
    unsigned short    type;
    int               (*func)();
    struct packet_type *next;
};
```

**Figure 14.7:** Structure identifying a packet type

```
struct iphdr{  
    u8  protocol;  
    u16 check;  
    u32 saddr;  
    u32 daddr;  
};
```

**Figure 14.8:** An IP header


```
struct udphdr{  
    u16 source;  
    u16 dest;  
    u16 check;  
};
```

**Figure 14.9:** A UDP header

```
struct tcphdr{  
    u16 source;  
    u16 dest;  
    u32 seq;  
    u16 check;  
};
```

**Figure 14.10:** A TCP header





**Figure 14.11:** Function call/return

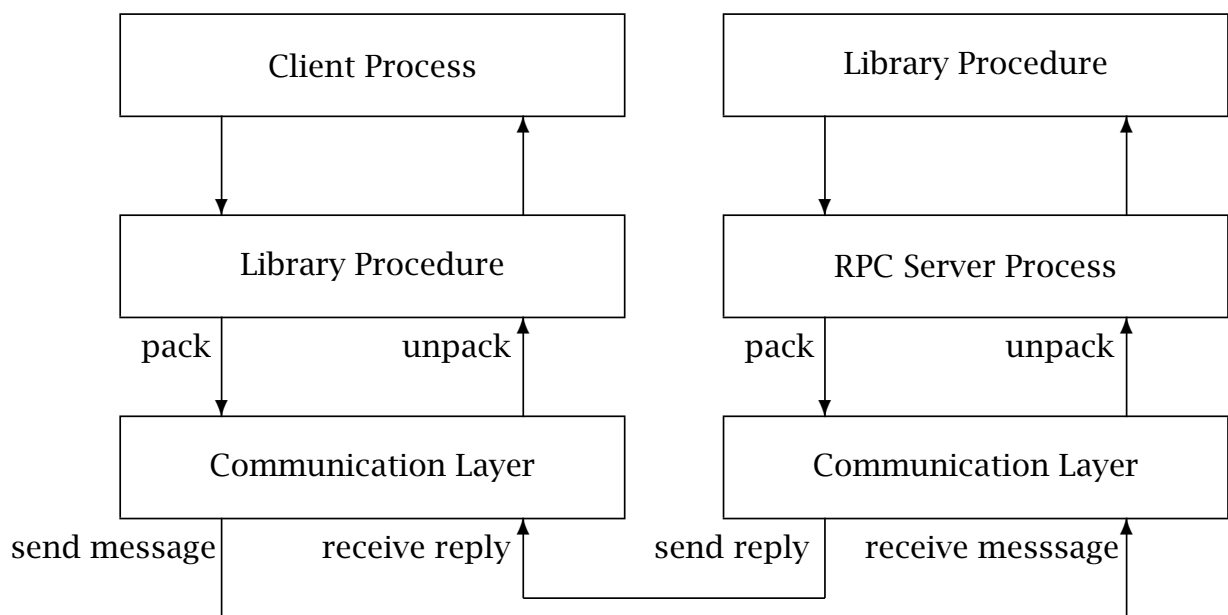


Figure 14.12: Call/return with RPC

```
program-definition:
    "program" program-name "{"
        version-list
    "}" "=" value ";"

version-list:
    version ";"
    | version ";" version list

version:
    "version" version-name "{"
        procedure-list
    "}" "=" value ";"

procedure-list:
    procedure ";"
    | procedure ";" procedure-list

procedure:
    type-ident procedure-name "(" type-ident ")" "=" value ";"
```

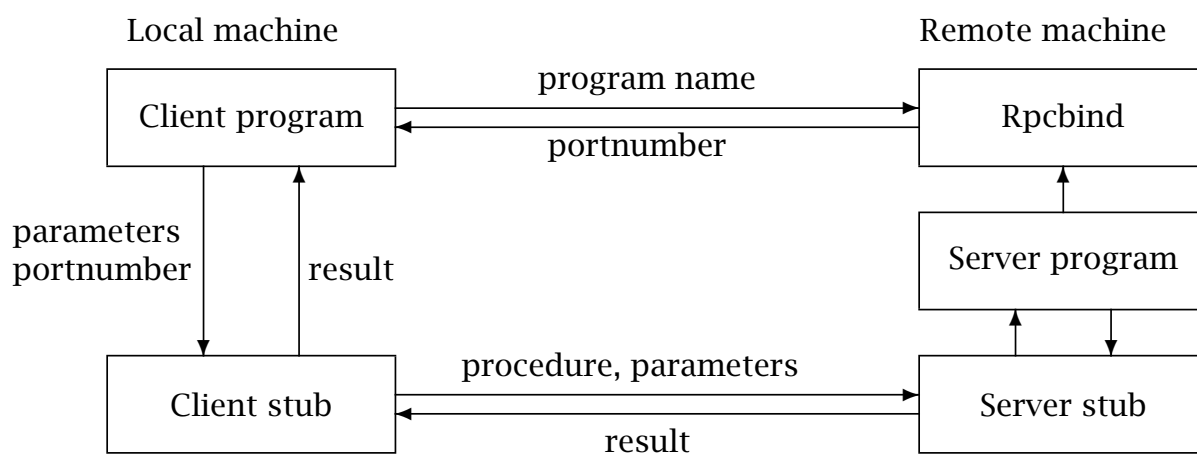
**Figure 14.13:** Format of RPC specification language

```
program MATH{  
    version MATHVERSION{  
        int DOUBLE (int) = 1;  
        } = 1;  
    } = 0x20000001;
```

**Figure 14.14:** Interface definition for MATH program

```
struct authunix_parms{
    u_long aup_time;      /* time credentials were created */
    char   *aup_machname; /* host name of client's machine */
    int     aup_uid;      /* client's Unix user id      */
    int     aup_gid;      /* client's current group id  */
    int     *aup_gids;    /* array of client's groups   */
};
```

**Figure 14.15:** Unix style credentials



**Figure 14.16:** Overview of RPC system

```
struct rpc_msg{
    u_long          rm_xid;
    enum msg_type   rm_direction;
    union{
        struct call_body  RM_cmb;
        struct reply_body RM_rmb;
    }ru;
};
```

**Figure 14.17:** Structure of an RPC message

```
struct call_body {  
    u_long      cb_rpcvers; /* RPC version number */  
    u_long      cb_prog;    /* program number    */  
    u_long      cb_vers;    /* version number */  
    u_long      cb_proc;    /* procedure number */  
    struct opaque_auth cb_cred; /* authentication */  
};
```

**Figure 14.18:** Body of a call message



```
struct reply_body{
    enum reply_stat      rp_stat;
    union{
        struct accepted_reply RP_ar;
        struct rejected_reply RP_dr;
    }ru;
};
```

**Figure 14.19:** Body of a reply message






Figure 15.1: Events in process X




Figure 15.2: Events in process Y




Figure 15.3: Events in process Z



**Figure 15.4:** Centralised control of mutual exclusion

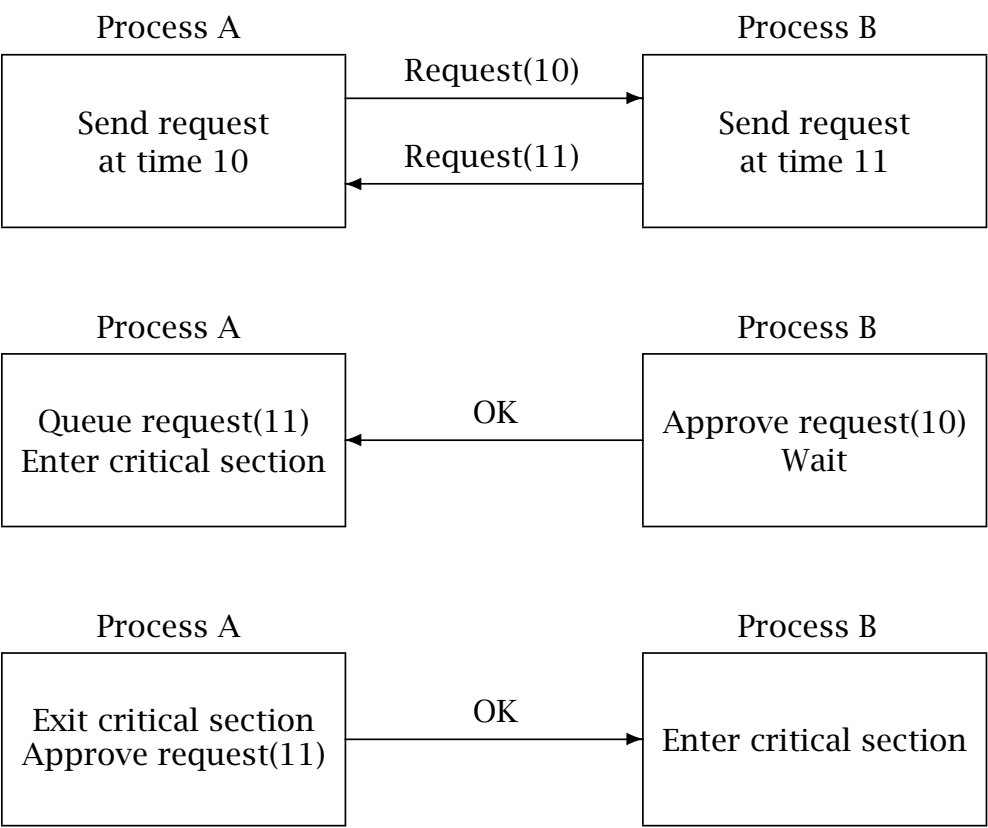
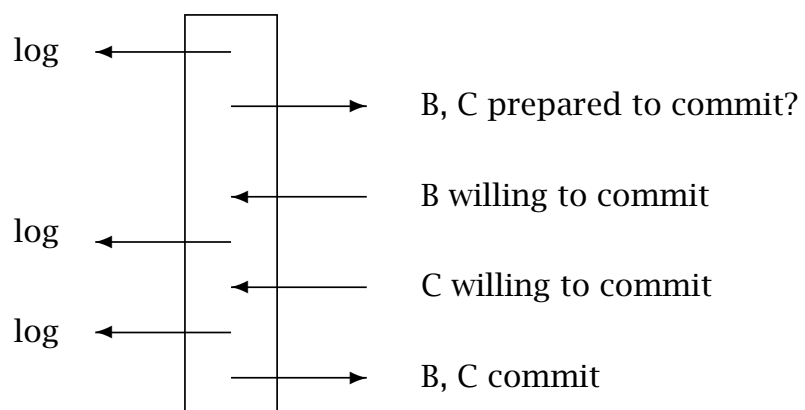
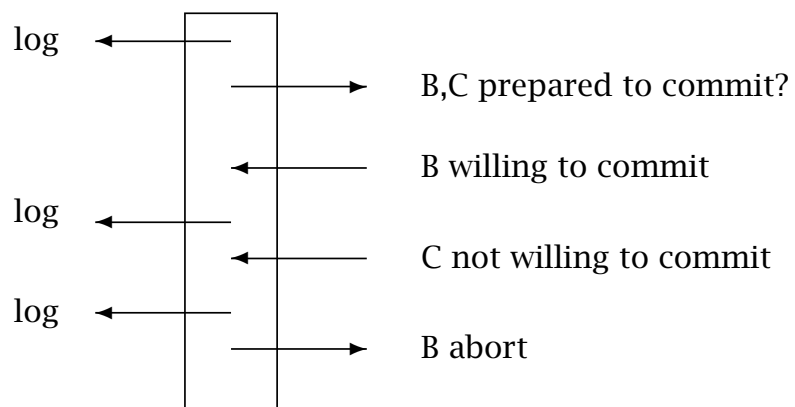


Figure 15.5: Distributed control of mutual exclusion

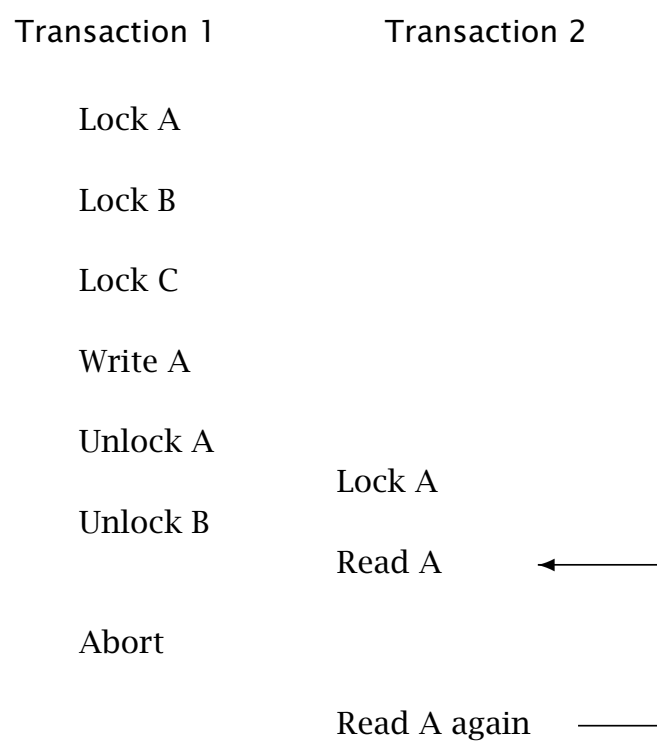


**Figure 15.6:** Successful two phase commit





**Figure 15.7:** Aborted two phase commit



**Figure 15.8:** A cascaded abort




Figure 15.9: A lock table

```
Lock (data-id, trans-id, lock-type)

    MUTEX_LOCK(table)
    IF (data-id already in table) THEN
        WHILE conflicting
            CWAIT(condvar-data-id)
        ENDWHILE
        Add trans-id to existing entry
    ENDIF
    Add new entry to table
    MUTEX_UNLOCK(table)
```

**Figure 15.10:** Algorithm for lock

```
Unlock(trans-id)

    MUTEX_LOCK(table)
    FOR (each entry in table belonging to trans-id) DO
        Remove the entry
        IF (trans-id was only holder) THEN
            CSIGNAL (condvar-data-id)
        ENDIF
    ENDFOR
    MUTEX_UNLOCK(table)
```

**Figure 15.11:** Algorithm for unlock






Figure 15.12: Local and global wait-for graphs



**Figure 16.1:** Situation before attaching remote files



**Figure 16.2:** Local file system after attaching remote files






Figure 16.3: A client and two servers




Figure 16.4: File system after mounting Dir4

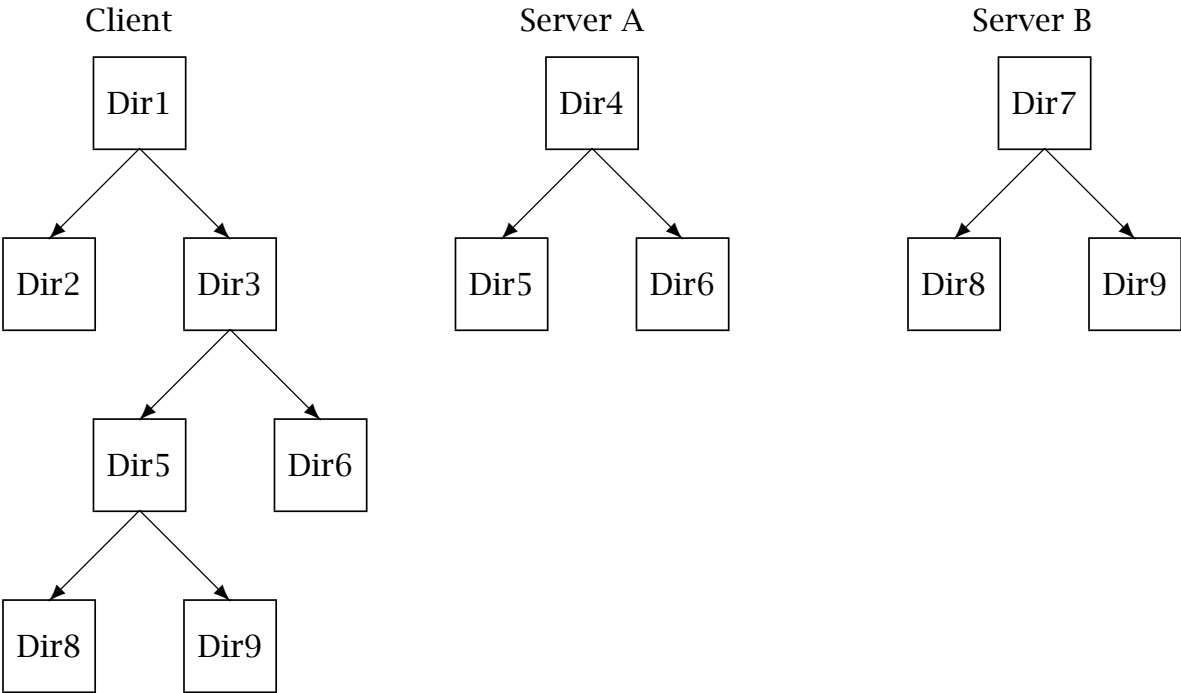
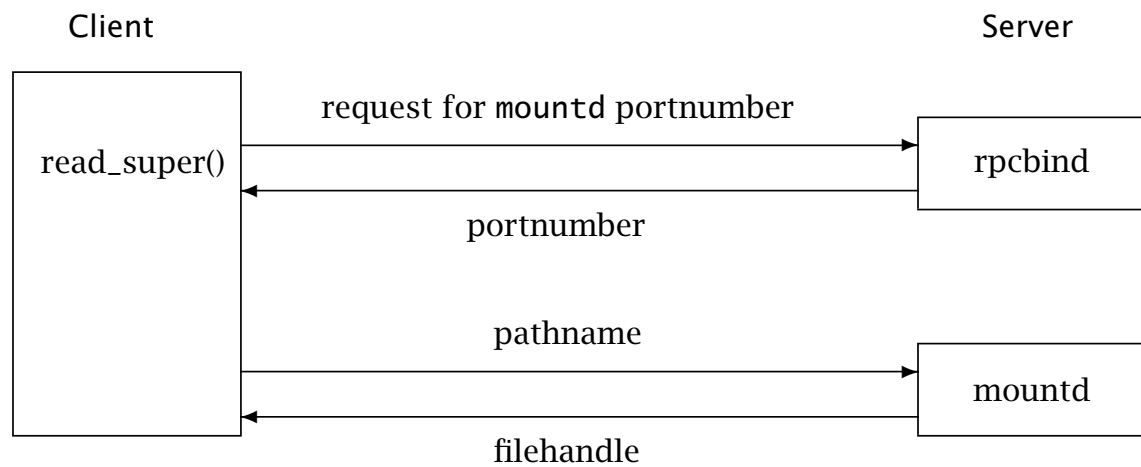


Figure 16.5: File system after mounting Dir7



**Figure 16.6:** Mounting a remote file system

```
struct nfs_sb_info{  
    struct nfs_server s_server;  
    struct nfs_fh     s_root;  
};
```

**Figure 16.7:** NFS specific super block information



**Figure 16.8:** Interaction between client and server

```
nfs_inode_info{  
    struct nfs_fh fhandle;  
    unsigned long read_cache_jiffies;  
    unsigned long read_cache_mtime;  
    unsigned long attrtimeo;  
};
```

**Figure 16.9:** NFS specific inode information

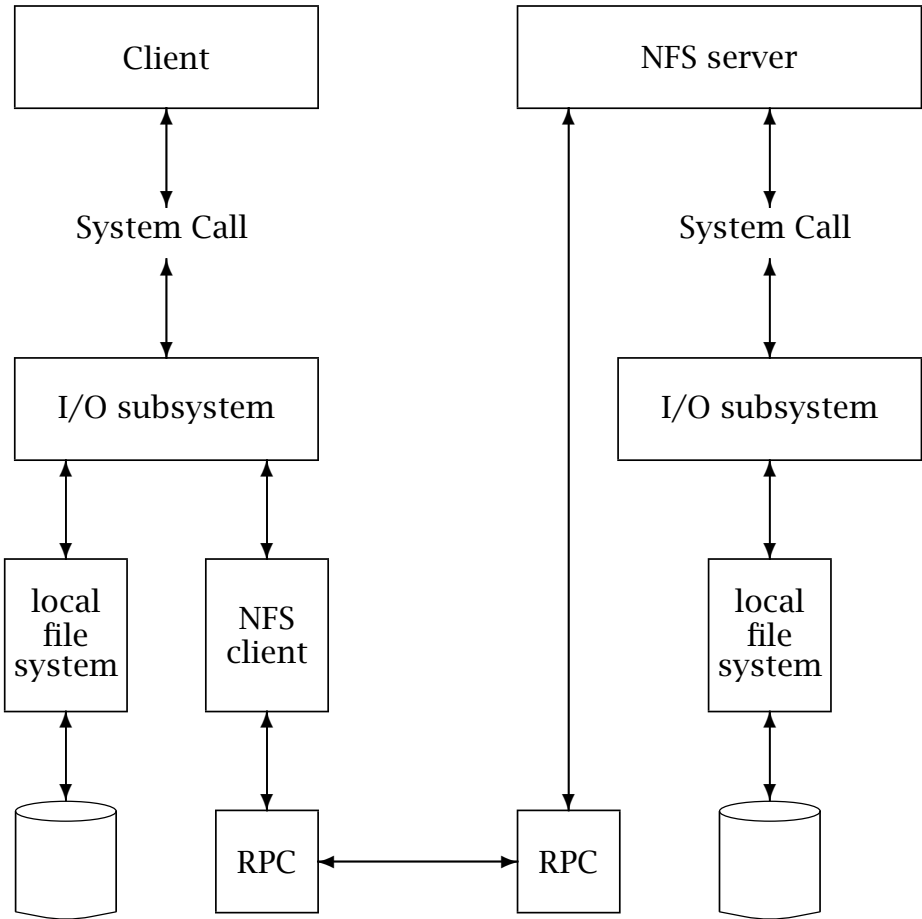


Figure 16.10: Handling local and remote file systems



	File1	File2	File3	CDdrive	Printer
Domain1	Read		Read		
Domain2				Read	Print
Domain3		Read	Execute		
Domain4	Read/Write		Read/Write		

**Figure 17.1:** Example access matrix

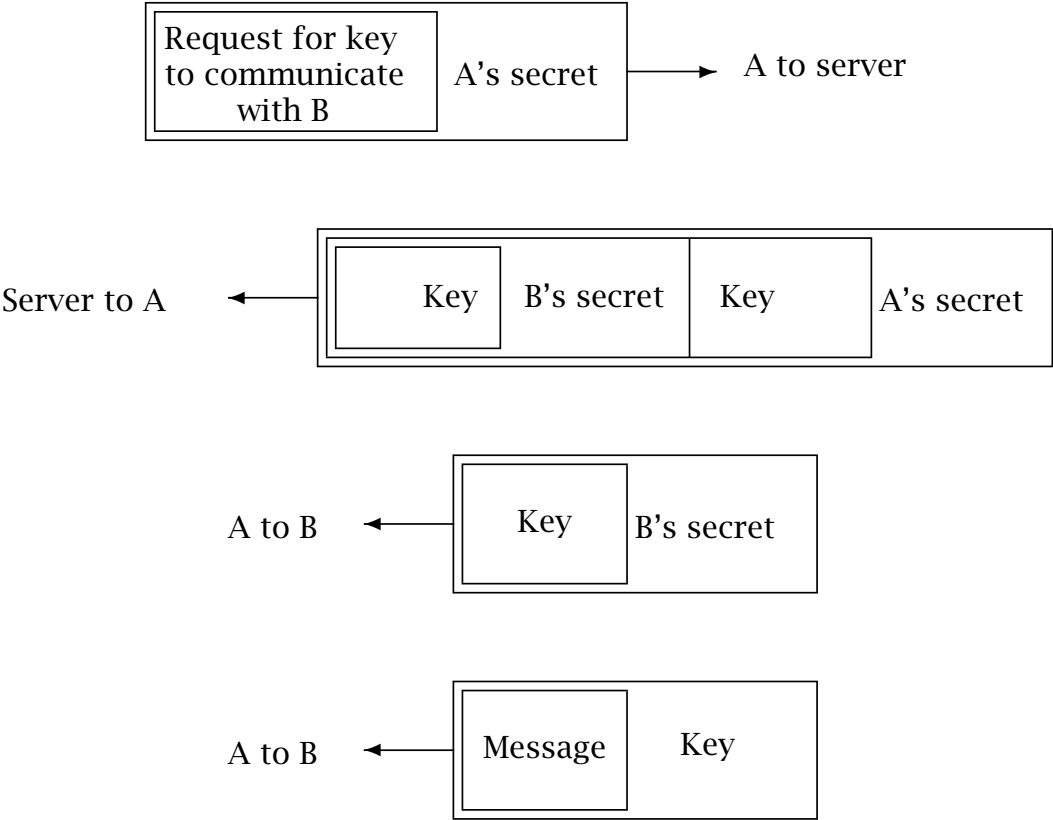


Figure 17.2: Sequence of messages with secret keys

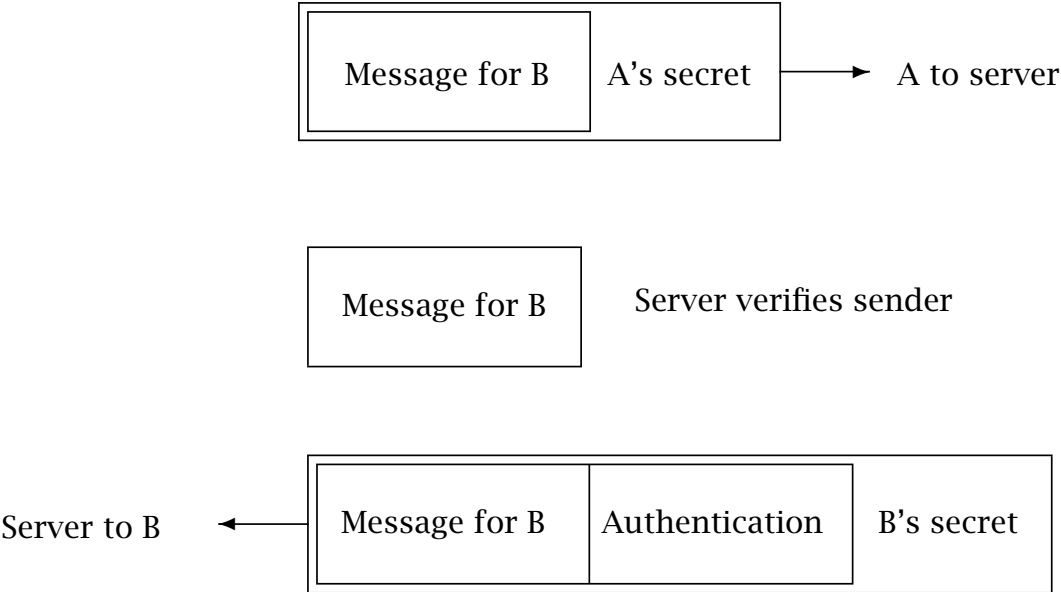
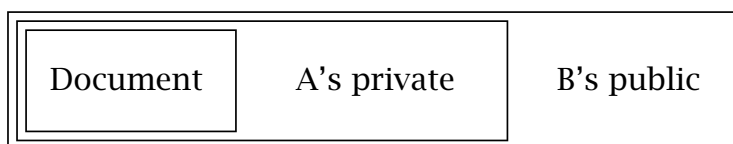



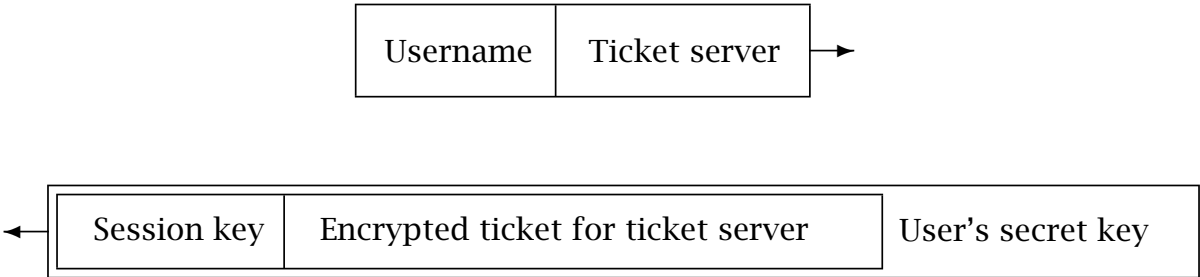
Figure 17.3: Digital signature with secret key




**Figure 17.4:** Encryption for both security and authentication




**Figure 17.5:** Three phases in the Kerberos system



**Figure 17.6:** Initial exchange of messages with authentication server



**Figure 17.7:** Authenticating a user to a file server



**Figure 17.8:** Authenticating a user to a file server





**Figure 17.9:** Requesting service of a file server